

Assignment2

11911839 聂雨荷

Q1

[20pts] `make qemu` 指令将指向 makefile 对应的 label, 该指令对应于:

```
1 qemu-system-riscv64 \  
2 -machine virt \  
3 -nographic \  
4 -bios default \  
5 -device loader,file=bin/ucore.bin,addr=0x80200000
```

请解释以上指令中每个参数的作用

- `-machine virt` : selects emulated machine `virt` (RISC-V VirtIO board)
- `-nographic` : disable graphical output and redirect serial I/Os to console
- `-bios default` : set the filename for the BIOS as default
- `-device loader,file=bin/ucore.bin,addr=0x80200000` : add misc device `loader` (call the bootloader OpenSBI), set the property `file` to be `bin/ucore.bin` and the qemu operation system starts at address `0x80200000`

Q2

[20pts] 请查阅资料, 理解并解释 `/lab/tools/kernel.ld` 文件以下片段中每一行的作用
(参考: <https://sourceware.org/binutils/docs/ld/Scripts.html>)

```
SECTIONS  
{  
    /* Load the kernel at this address: "." means the current address */  
    . = BASE_ADDRESS;  
  
    .text : {  
        *(.text.kern_entry)  
        *(.text.stub .text.* .gnu.linkonce.t.*)  
    }  
  
    PROVIDE(etext = .); /* Define the 'etext' symbol to this value */  
  
    .rodata : {  
        *(.rodata .rodata.* .gnu.linkonce.r.*)  
    }  
  
    /* Adjust the address for the data segment to the next page */  
    . = ALIGN(0x1000);
```

```

1  SECTIONS{
2      ...
3  }

```

- `SECTION` command describes the memory layout of the output file

```

1  /* Load the kernel at this address: `` means the current address */
2  . = BASE_ADDRESS;

```

- set the value of the location counter `.` to be `BASE_ADDRESS`

```

1  .text : {
2      *(.text.kern_entry)
3      *(.text .stub .text.* .gnu.linkonce.t.*)
4  }

```

- define an output section `.text` (program code)
- `*(.text.kern_entry)` : add any input sections ended by `.text.kern_entry` in all input files
- `*(.text .stub .text.* .gnu.linkonce.t.*)` : add any input sections ended by `.text` / `.stub` / `.text.*` / `.gnu.linkonce.t.*`
- since the location counter is `BASE_ADDRESS`, when the output section `.text` is defined, the linker will set the address of the `.text` section in the output file to be `BASE_ADDRESS`.

```

1  /* Define the 'etext' symbol to this value */
2  PROVIDE(etext = .);

```

- define an `etext` symbol, and set the value of the symbol to be the location counter `.`

```

1  .rodata : {
2      *(.rodata .rodata.* .gnu.linkonce.r.*)
3  }

```

- define an output section `.rodata` (read-only data)
- `*(.rodata .rodata.* .gnu.linkonce.r.*)` : add any input sections ended by `.rodata` / `.rodata.*` / `.gnu.linkonce.r.*`
- since the location counter doesn't change, the `.rodata` output section is allocated directly after the `.text` output section

```

1  /* Adjust the address for the data segment to the next page */
2      . = ALIGN(0x1000);

```

- adjust the address for the data segment to the next page

Q3

[10pts] 请解释 `/lab/kern/init/init.c` 中 `main` 函数中 `memset(edata, 0, end - edata);` 的参数及语句作用。（需要读到的代码有 `init.c`, `kernel.ld`）

- `edata` : the pointer to the start block of the memory to be filled
- `0` : memory address to be set
- `end - edata` : number of bytes start from the `edata` to be set
- `edata` 和 `end` 分别记录了链接器输出的 `.bss` (read-write zero initialized data.) 的起始位置和终止位置的地址
- `memset()` 函数将 `end-edata` 个 byte, `edata` 中的也就是 `.bss` 文件的数据放入 memory 的地址 `0` 处, 完成初始化

Q4

[20pts] 请描述 `cputs()` 指令是如何通过 `sbi` 打印字符的。

函数顺序如下:

```

1  /* *
2   * cputs- writes the string pointed by @str to stdout and
3   * appends a newline character.
4   * */
5  int cputs(const char *str) {
6      int cnt = 0;
7      char c;
8      while ((c = *str++) != '\0') {
9          cputch(c, &cnt);
10     }
11     cputch('\n', &cnt);
12     return cnt;
13 }
14
15 /* *
16 * cputch - writes a single character @c to stdout, and it will
17 * increace the value of counter pointed by @cnt.
18 * */
19 static void cputch(int c, int *cnt) {
20     cons_putc(c);

```

```

21     (*cnt)++;
22 }
23
24 /* cons_putc - print a single character @c to console devices */
25 void cons_putc(int c) {
26     sbi_console_putchar((unsigned char)c);
27 }
28
29 void sbi_console_putchar(unsigned char ch) {
30     sbi_call(SBI_CONSOLE_PUTCHAR, ch, 0, 0);
31 }
32
33 uint64_t sbi_call(uint64_t sbi_type, uint64_t arg0, uint64_t arg1, uint64_t
34 arg2) {
35     uint64_t ret_val;
36     __asm__ volatile (
37         "mv x17, %[sbi_type]\n"
38         "mv x10, %[arg0]\n"
39         "mv x11, %[arg1]\n"
40         "mv x12, %[arg2]\n" //mv操作把参数的数值放到寄存器里
41         "ecall\n" //参数放好之后, 通过ecall, 交给OpenSBI来执行
42         "mv %[ret_val], x10"
43         //OpenSBI按照riscv的calling convention,把返回值放到x10寄存器里
44         //我们还需要自己通过内联汇编把返回值拿到我们的变量里
45         : [ret_val] "=r" (ret_val)
46         : [sbi_type] "r" (sbi_type), [arg0] "r" (arg0), [arg1] "r" (arg1),
47           [arg2] "r" (arg2)
48         : "memory"
49     );
50     return ret_val;
51 }
52

```

- 传入字符串指针 `*str`，每一个字符串以 `\0` 作为结尾
- 遍历字符串中每一个字符，对于每一个字符，调用 `cputch()`，将单个字符写入stdout，它将增加由 `cnt` 指向的 `counter` 的值
- `cputch()` 是对 `sbi_console_putchar()` 的简单封装，字符继续向下传入
- `sbi_console_putchar()` 向下调用 `sbi_ecall()`，向 OpenSBI 传入 `ecall` 的指令是 `SBI_CONSOLE_PUTCHAR`，并且传递字符
- `sbi_call()` 使用内联汇编，将包括 `sbi_type` 和字符放对应位置后，调用 `ecall` 指令，启动调用 OpenSBI 进入 M 态处理

Q5

[30pts] 编程题 请在第三周 lab.zip 代码包的基础上，理解使用 `ecall` 打印字符的原理，实现一个 `shutdown()` 关机函数。（所有修改到的代码请截图和运行结果截图一起放在报告中）

在 `while(1);` 前面填加：

```

1 //init.c
2 //-----
3 cputs("The system will close.\n");
4 shutdown();
5 // -----
6 while (1)
7     ;

```

实现效果(代码不会执行到while语句):

参考资料: <https://github.com/riscv-non-isa/riscv-sbi-doc/blob/master/riscv-sbi.adoc#410-function-listing>

```

lab > libs > C sbi.c > ...
19     "mv x10, %[arg0]\n"
20     "mv x11, %[arg1]\n"
21     "mv x12, %[arg2]\n"
22     "ecall\n"
23     "mv %[ret_val], x10"
24     : [ret_val] "=r" (ret_val)
25     : [sbi_type] "r" (sbi_type), [arg0] "r" (arg0), [arg1] "r" (arg1)
26     : "memory"
27 );
28 return ret_val;
29 }
30
31 int sbi_console_getchar(void) {
32     return sbi_call(SBI_CONSOLE_GETCHAR, 0, 0, 0);
33 }
34 void sbi_console_putchar(unsigned char ch) {
35     sbi_call(SBI_CONSOLE_PUTCHAR, ch, 0, 0);
36 }
37
38 void sbi_set_timer(unsigned long long stime_value) {
39     sbi_call(SBI_SET_TIMER, stime_value, 0, 0);
40 }
41
42 /*
43 [Assignment2] SBI Shutdown
44 */
45 void sbi_shutdown(void){
46     sbi_call(SBI_SHUTDOWN, 0, 0, 0);
47 }
48

```

EXPLORER

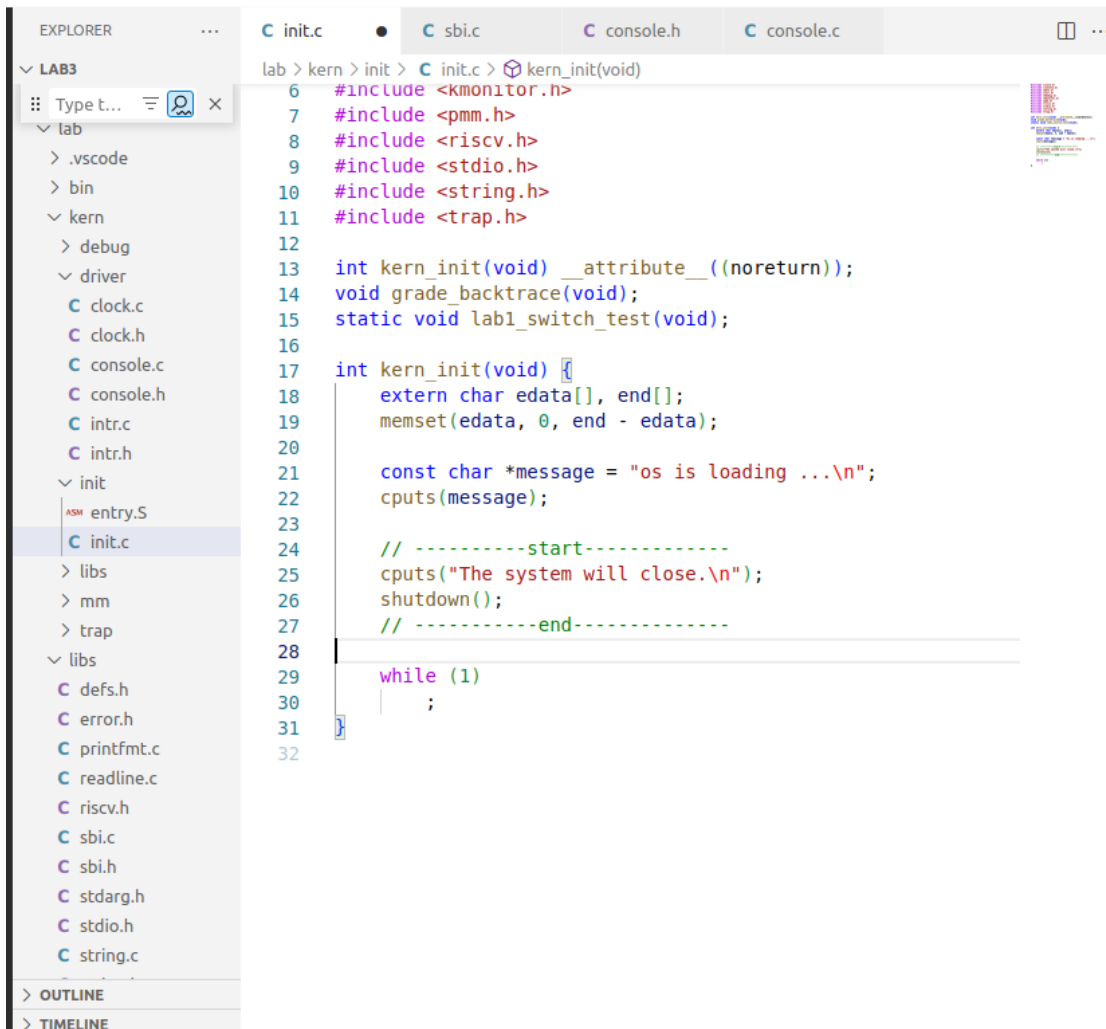
lab > kern > driver > console.h > shutdown(void)

```
1  #ifndef __KERN_DRIVER_CONSOLE_H
2  #define __KERN_DRIVER_CONSOLE_H
3
4  void cons_init(void);
5  void cons_putc(int c);
6  int cons_getc(void);
7  void serial_intr(void);
8  void kbd_intr(void);
9  void shutdown(void); // Assignment2
10
11 #endif /* !__KERN_DRIVER_CONSOLE_H */
12
13
```

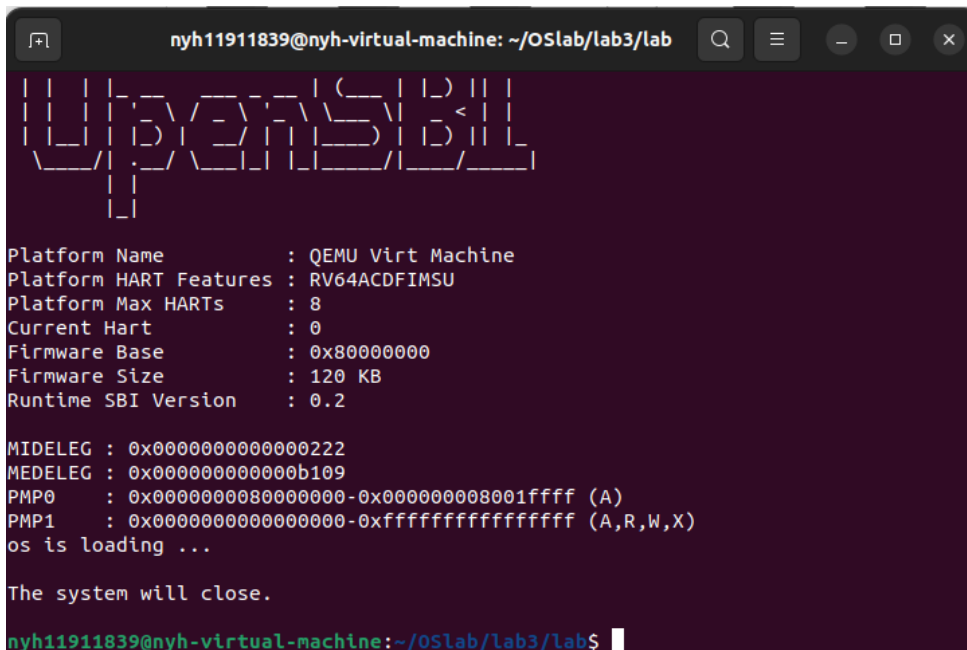
EXPLORER

lab > kern > driver > console.c > ...

```
6
7  /* serial_intr - try to feed input characters from serial port */
8  void serial_intr(void) {}
9
10 /* cons_init - initializes the console devices */
11 void cons_init(void) {}
12
13 /* cons_putc - print a single character @c to console devices */
14 void cons_putc(int c) { sbi_console_putchar((unsigned char)c); }
15
16 /* *
17  * cons_getc - return the next input character from console,
18  * or 0 if none waiting.
19  * */
20 int cons_getc(void) {
21     int c = 0;
22     c = sbi_console_getchar();
23     return c;
24 }
25
26 /*
27 [Assignment2] shutdown
28 */
29 void shutdown(void) { sbi_shutdown(); }
30
```



```
lab > kern > init > C init.c > kern_init(void)
6  #include <kmonitor.h>
7  #include <pmm.h>
8  #include <riscv.h>
9  #include <stdio.h>
10 #include <string.h>
11 #include <trap.h>
12
13 int kern_init(void) __attribute__((noreturn));
14 void grade_backtrace(void);
15 static void lab1_switch_test(void);
16
17 int kern_init(void) {
18     extern char edata[], end[];
19     memset(edata, 0, end - edata);
20
21     const char *message = "os is loading ...\n";
22     cputs(message);
23
24     // -----start-----
25     cputs("The system will close.\n");
26     shutdown();
27     // -----end-----
28
29     while (1)
30     ;
31 }
32
```



```
nyh11911839@nyh-virtual-machine: ~/OSlab/lab3/lab
Up and Running!

Platform Name       : QEMU Virt Machine
Platform HART Features : RV64ACDFIMSU
Platform Max HARTs   : 8
Current Hart        : 0
Firmware Base       : 0x80000000
Firmware Size       : 120 KB
Runtime SBI Version  : 0.2

MIDELEG : 0x0000000000000222
MEDELEG : 0x000000000000b109
PMP0    : 0x0000000080000000-0x000000008001ffff (A)
PMP1    : 0x0000000000000000-0xffffffffffff (A,R,W,X)
os is loading ...

The system will close.
nyh11911839@nyh-virtual-machine:~/OSlab/lab3/lab$
```