

Week5 Report in Class (Fri56)

11911839 聂雨荷

Q1

代码中如何区分父子进程?父子进程的执行顺序是否是固定的?

- We can use the return value of `fork()` function to determine which of process is the parent and which is the child, on success if
 - `fork() == 0` , current process is the child process
 - `fork() == childPID` , current process is the parent process
- The execution order of parent-child processes is not fixed

Q2

请回答第四步僵尸进程中列举的第4种情况的结果会是什么。

- If the child process does not exit but the parent process does exit, it will become orphan process (孤儿进程).
- The child process is adopted to the `init` process (PID=1) or to the registered grandfather process.

Q3

请编写一段c语言代码 (截图), 用于产生僵尸进程, 并截图僵尸进程的状态 (ps).

```
1  #include <stdio.h>
2  #include <unistd.h>
3  #include <errno.h>
4  #include <stdlib.h>
5
6  int main()
7  {
8      pid_t pid = fork();
9      if (pid < 0)
10     {
11         perror("fork error:");
12         exit(1);
13     }
14     else if (pid == 0)
```

```

15     {
16         printf("[%d] I am child process. Byebye.\n", getpid());
17         exit(0);
18     }
19     printf("[%d] I am parent process. I will sleep two seconds\n", getpid());
20     sleep(2);
21     system("ps -o pid,ppid,state,TTY,command");
22     printf("parent process is exiting.\n");
23     return 0;
24 }

```

```

nyh11911839@nyh-virtual-machine: ~/OSlab/lab5/lab5-q
nyh11911839@nyh-virtual-machine:~/OSlab/lab5/lab5-q$ gcc q3.c
nyh11911839@nyh-virtual-machine:~/OSlab/lab5/lab5-q$ ./a.out
[8276] I am parent process. I will sleep two seconds
[8277] I am child process. Byebye.
  PID  PPID S  TT  COMMAND
 8143   8125 S  pts/3  bash
 8276   8143 S  pts/3  ./a.out
 8277   8276 Z  pts/3  [a.out] <defunct>
 8278   8276 S  pts/3  sh -c ps -o pid,ppid,state,TTY,command
 8279   8278 R  pts/3  ps -o pid,ppid,state,TTY,command
parent process is exiting.
nyh11911839@nyh-virtual-machine:~/OSlab/lab5/lab5-q$

```

Q4

lab5 的 ucore 代码具体通过哪条指令以什么形式跳转至 init_main()

- `init_main` 不断调用 `cpu_idle()` 函数，进而调用到 `schedule()` 函数，`schedule` 会找到允许调度的进程
- 我们会在这中间调用 `forkets` 函数将 `init_main` 进程的中断帧放在 `sp`，从中恢复所有的寄存器，通过 `kernel_thread_entry`，我们将 `s0` 寄存器存放着新进程执行的函数，`s1` 存放着传给函数的参数，我们把参数放在 `a0` 寄存器，并跳转到 `s0` 执行 `init_main`，从而完成调用

```

1  .text
2  .globl kernel_thread_entry
3  kernel_thread_entry:
4      move a0, s1
5      jalr s0 # jump to init_main() function
6      jal do_exit

```

Q5

lab5 的 ucore 代码中是如何调用到 kernel_thread_entry 的?

1. initial `init_main` process, idle process
2. call `kernel_thread` function to create a new process `init_main`, and it will add this process to the process linked list, below 3-5 explain the details
3. `kernel_thread` will call `do_fork` function to initial a new PCB

```
1 // kernel_thread - create a kernel thread using "fn" function
2 // NOTE: the contents of temp trapframe tf will be copied to
3 //       proc->tf in do_fork-->copy_thread function
4 int
5 kernel_thread(int (*fn)(void *), void *arg, uint32_t clone_flags) {
6     struct trapframe tf;
7     memset(&tf, 0, sizeof(struct trapframe));
8
9     tf.gpr.s0 = (uintptr_t)fn;
10    tf.gpr.s1 = (uintptr_t)arg;
11    tf.status = (read_csr(sstatus) | SSTATUS_SPP | SSTATUS_SPIE) & ~SSTATUS_SIE;
12    tf.epc = (uintptr_t)kernel_thread_entry;
13    return do_fork(clone_flags | CLONE_VM, 0, &tf);
14 }
```

4. `do_fork` will call the `copy_thread` function to copy the current process, besides, it will add the new process to the process linked list

```
1 /* do_fork -      parent process for a new child process
2  * @clone_flags: used to guide how to clone the child process
3  * @stack:       the parent's user stack pointer. if stack==0, It means to fork a
4  *               kernel thread.
5  * @tf:          the trapframe info, which will be copied to child process's
6  *               proc->tf
7  */
8 int
9 do_fork(uint32_t clone_flags, uintptr_t stack, struct trapframe *tf) {
10     int ret = -E_NO_FREE_PROC;
11     struct proc_struct *proc;
12     if (nr_process >= MAX_PROCESS) {
13         goto fork_out;
14     }
15     ret = -E_NO_MEM;
16     // 1. call alloc_proc to allocate a proc_struct
17     // 2. call setup_kstack to allocate a kernel stack for child process
18     // 3. call copy_mm to dup OR share mm according clone_flag
19     // 4. call copy_thread to setup tf & context in proc_struct
```

```

18     // 5. insert proc_struct into hash_list && proc_list
19     // 6. call wakeup_proc to make the new child process RUNNABLE
20     // 7. set ret vaule using child proc's pid
21     if ((proc = alloc_proc()) == NULL) {
22         goto fork_out;
23     }
24     if ((ret = setup_kstack(proc)) == -E_NO_MEM) {
25         goto bad_fork_cleanup_proc;
26     }
27     copy_mm(clone_flags, proc);
28
29     copy_thread(proc, stack, tf);
30
31     const int pid = get_pid();
32     proc->pid = pid;
33     list_add(hash_list + pid_hashfn(pid), &(proc->hash_link));
34     list_add(&proc_list, &(proc->list_link));
35     nr_process++;
36
37     wakeup_proc(proc);
38     ret = pid;
39 fork_out:
40     return ret;
41
42 bad_fork_cleanup_kstack:
43     put_kstack(proc);
44 bad_fork_cleanup_proc:
45     kfree(proc);
46     goto fork_out;
47 }

```

5. `copy_thread` will copy and switch a new process, besides, it sets the `ra` to be the entry of `forkret` function

```

1 // copy_thread - setup the trapframe on the process's kernel stack top and
2 //               - setup the kernel entry point and stack of process
3 static void
4 copy_thread(struct proc_struct *proc, uintptr_t esp, struct trapframe *tf) {
5     proc->tf = (struct trapframe *) (proc->kstack + KSTACKSIZE - sizeof(struct
6     trapframe));
7
8     *(proc->tf) = *tf;
9
10    // Set a0 to 0 so a child process knows it's just forked
11    proc->tf->gpr.a0 = 0;
12    proc->tf->gpr.sp = (esp == 0) ? (uintptr_t)proc->tf : esp;
13
14    proc->context.ra = (uintptr_t)forkret;
15    proc->context.sp = (uintptr_t)(proc->tf);

```

```
14 }
```

6. since `idle_proc` need reschedule, it will call `schedule` function to check whether there is a process schedulable

```
1 void cpu_idle(void) {
2     while (1) {
3         if (current->need_resched) {
4             schedule();
5         }
6     }
7 }
8
```

7. `schedule` will call `proc_run` function to wake up selected process

```
1 void
2 schedule(void) {
3     bool intr_flag;
4     list_entry_t *le, *last;
5     struct proc_struct *next = NULL;
6     local_intr_save(intr_flag);
7     {
8         current->need_resched = 0;
9         last = (current == idleproc) ? &proc_list : &(current->list_link);
10        le = last;
11        do {
12            if ((le = list_next(le)) != &proc_list) {
13                next = le2proc(le, list_link);
14                if (next->state == PROC_RUNNABLE) {
15                    break;
16                }
17            }
18        } while (le != last);
19        if (next == NULL || next->state != PROC_RUNNABLE) {
20            next = idleproc;
21        }
22        next->runs ++;
23        if (next != current) {
24            proc_run(next);
25        }
26    }
27    local_intr_restore(intr_flag);
28 }
```

8. `proc_run` will call `switch_to` to switch new process

```
1 // proc_run - make process "proc" running on cpu
2 // NOTE: before call switch_to, should load base addr of "proc"'s new PDT
3 void
4 proc_run(struct proc_struct *proc) {
5     if (proc != current) {
6         bool intr_flag;
7         struct proc_struct *prev = current, *next = proc;
8         local_intr_save(intr_flag);
9         {
10             current = proc;
11             lcr3(next->cr3);
12             switch_to(&(prev->context), &(next->context));
13         }
14         local_intr_restore(intr_flag);
15     }
16 }
```

9. `switch_to` function save and swap registers that need to be saved, `ra` register is set to be the entry of `forkret` function, so it's going to return the `forkret` function

```
1 #include <riscv.h>
2
3 .text
4 # void switch_to(struct proc_struct* from, struct proc_struct* to)
5 .globl switch_to
6 switch_to:
7     # save from's registers
8     STORE ra, 0*REGBYTES(a0) # here proc->context.ra = (uintptr_t)forkret;
9     STORE sp, 1*REGBYTES(a0)
10    STORE s0, 2*REGBYTES(a0)
11    STORE s1, 3*REGBYTES(a0)
12    STORE s2, 4*REGBYTES(a0)
13    STORE s3, 5*REGBYTES(a0)
14    STORE s4, 6*REGBYTES(a0)
15    STORE s5, 7*REGBYTES(a0)
16    STORE s6, 8*REGBYTES(a0)
17    STORE s7, 9*REGBYTES(a0)
18    STORE s8, 10*REGBYTES(a0)
19    STORE s9, 11*REGBYTES(a0)
20    STORE s10, 12*REGBYTES(a0)
21    STORE s11, 13*REGBYTES(a0)
22
23    # restore to's registers
24    LOAD ra, 0*REGBYTES(a1)
25    LOAD sp, 1*REGBYTES(a1)
```

```

26     LOAD s0, 2*REGBYTES(a1)
27     LOAD s1, 3*REGBYTES(a1)
28     LOAD s2, 4*REGBYTES(a1)
29     LOAD s3, 5*REGBYTES(a1)
30     LOAD s4, 6*REGBYTES(a1)
31     LOAD s5, 7*REGBYTES(a1)
32     LOAD s6, 8*REGBYTES(a1)
33     LOAD s7, 9*REGBYTES(a1)
34     LOAD s8, 10*REGBYTES(a1)
35     LOAD s9, 11*REGBYTES(a1)
36     LOAD s10, 12*REGBYTES(a1)
37     LOAD s11, 13*REGBYTES(a1)
38
39     ret # jump to the ra register
40

```

10. `forkret` function will enter `forkrets` function. This ends up how the program is calling those two functions.

```

1  // forkret -- the first kernel entry point of a new thread/process
2  // NOTE: the addr of forkret is setted in copy_thread function
3  //      after switch_to, the current proc will execute here.
4  static void
5  forkret(void) {
6      forkrets(current->tf);
7  }

```

11. `forkrets` function sets stack to this new process's trapframe and `__trapret` restores all the registers directly from inside the interrupt frame

```

1      .globl forkrets
2  forkrets:
3      # set stack to this new process's trapframe
4      move sp, a0
5      j __trapret

```

12. the trapframe contains `epc` register which points to the `kernel_thread_entry`, so the program jumps to where the `kernel_thread_entry` function is located

```
1  .text
2  .globl kernel_thread_entry
3  kernel_thread_entry: # void kernel_thread(void)
4      move a0, s1
5      jalr s0
6      jal do_exit
```