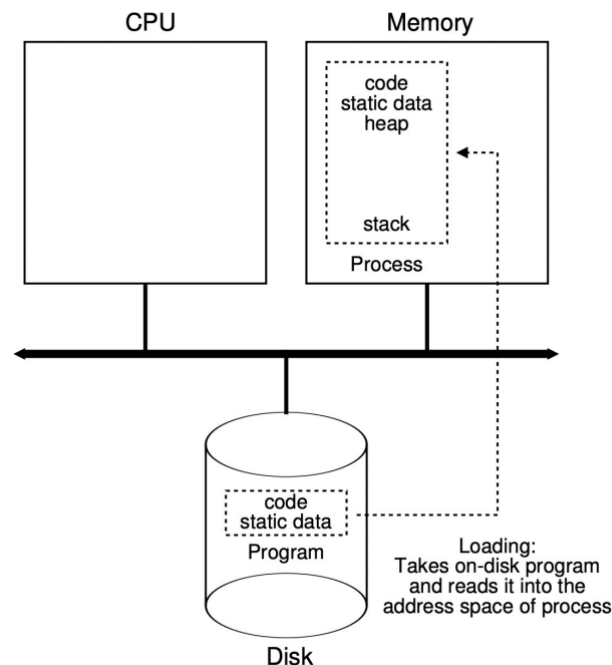



Lecture3 Processes

1. Process and System calls

What is a Process



- Process is a **program in execution**
- Program is a **file on the disk** (code and static data)
- Process is **loaded by the OS**
 - Code and static data are loaded from the program
 - Heap and stack are created by the OS



- A process is an abstraction of **machine states**
 - **Memory:** address space

- **Register:**
 - Program Counter (PC) or Instruction Pointer
 - Stack pointer
 - Frame pointer
- **I/O:** all files opened by the process

Process Identification

```
// compile to getpid
#include <stdio.h> // printf()
#include <unistd.h> // getpid()


int main(void) {
    printf("My PID is %d\n", getpid());
}
```

```
$ ./getpid
My PID is 1234
$ ./getpid
My PID is 1235
$ ./getpid
My PID is 1237
```

Each process is given a unique ID number, and is called the process ID, or the PID.

- System call `getpid()` prints the PID of the calling process

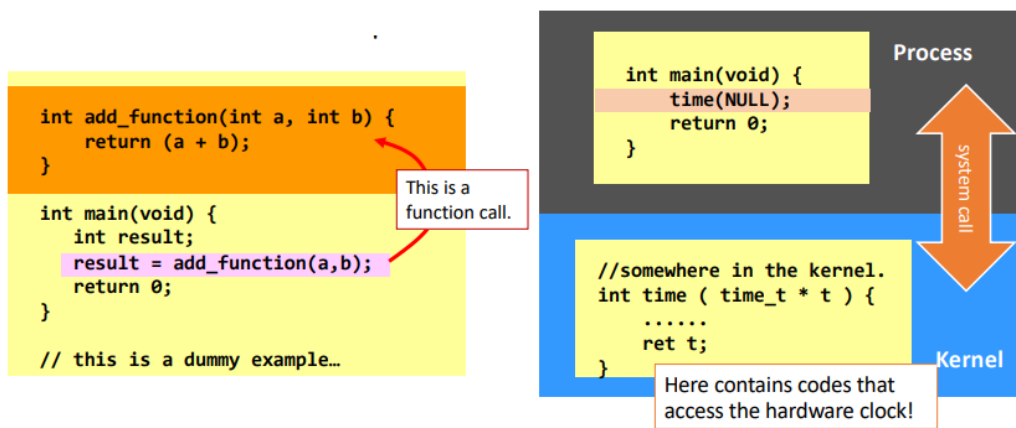
Process Life Cycle



- **interrupt** e.g. time interrupt
- **I/O or event wait** e.g. open a file

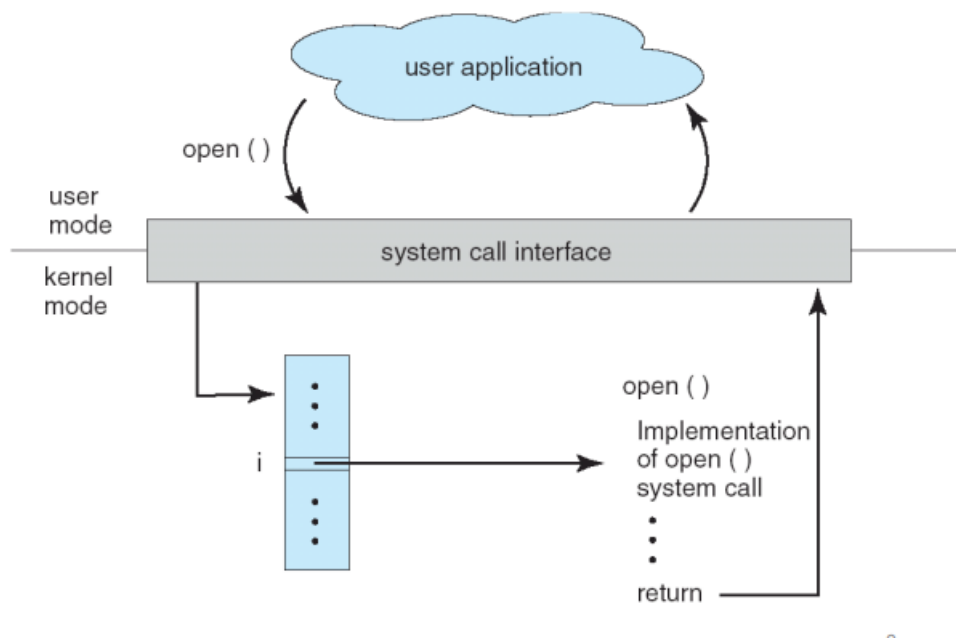
System Call

Process-Kernel Interaction



- System call is a **function call**
 - exposed by the **kernel**
 - **abstraction** of kernel operations

Call by Number



- System call is different from function call
- System call is a call **by number**

- User-mode code from xv6-riscv

```
int main(void) {
    .....
    int fd = open("copyin1", O_CREATE|O_WRONLY);
    .....
    return 0;
}
```

```
/* kernel/syscall.h */
```

```
#define SYS_open 15
```

```
/* user/usys.S */
.global open
open:
    li a7, SYS_open
    ecall
    ret
```

- Kernel code from xv6-riscv

```
/* kernel/syscall.h */
```

```
#define SYS_open 15
```

```
/* kernel/file.c */
```

```
uint64 sys_open(void) {
    .....
    return fd;
}
```

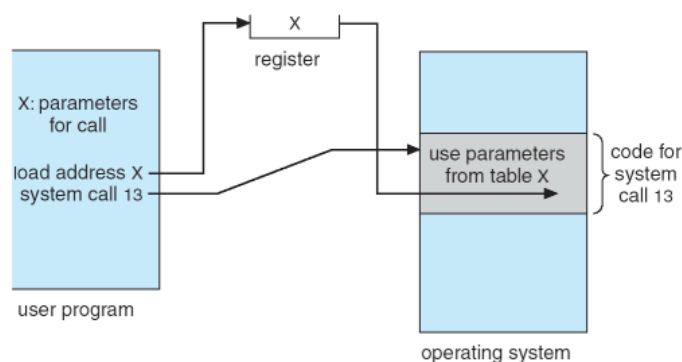
```
/* kernel/syscall.c */
```

```
static uint64 (*syscalls[])(void) = {
    .....
    [SYS_open] sys_open,
    .....
}
```

```
void syscall(void) {
    struct proc *p = myproc();
    num = p->trapframe->a7;
    p->trapframe->a0 = syscalls[num]();
}
```


Parameter Passing

- **more information** is required than the index of desired system call
 - Exact type and amount of information **vary** according to OS and call
- **Three general methods** used to pass parameters to the OS
 - **Registers:** pass the parameters in registers
 - In some cases, may be more parameters than registers
 - **x86** and **risc-v** take this approach
 - **Blocks:** Parameters stored in a memory block and address of the block passed as a parameter in a register



- **Stack:** Parameters placed, or pushed, onto the stack by the program and popped off the stack by the operating system
- Block and stack methods **do not limit the number or length** of parameters being passed

System Call v.s. Library API call

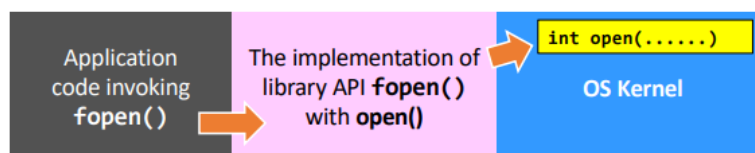


- Most operating systems provide **standard C library** to provide library API calls
 - A layer of indirection for system calls

Name	System call?
<code>printf()</code> & <code>scanf()</code>	No
<code>malloc()</code> & <code>free()</code>	No
<code>fopen()</code> & <code>fclose()</code>	No
<code>mkdir()</code> & <code>rmdir()</code>	Yes
<code>chown()</code> & <code>chmod()</code>	Yes




Library call	<code>fopen("hello.txt", "w");</code>
System call	<code>open("hello.txt", O_WRONLY O_CREAT O_TRUNC, 0666);</code>



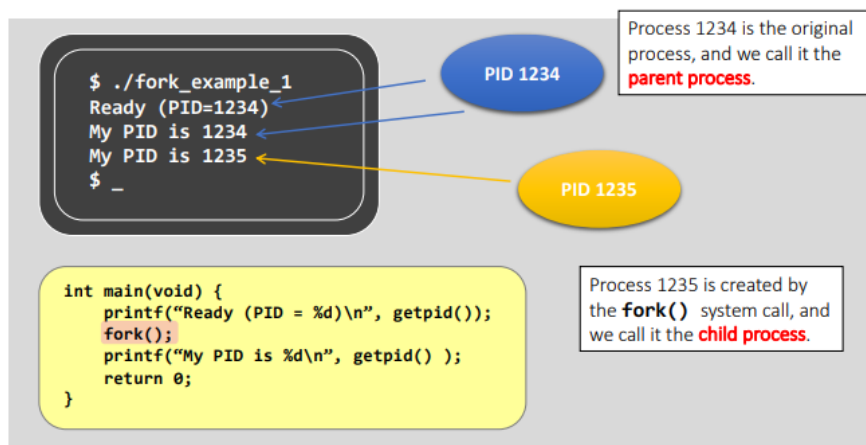
- Take `fopen()` as an example
 - `fopen()` invokes the system call `open()`
 - `open()` is too primitive and is not programmer-friendly

2. Process creation



- Parent process create children processes, which, in turn create other processes, forming a **tree of processes**. We create one process (parent process) from another process (child process).
- Generally, process identified and managed via a **process identifier** (PID)

Creating Process with fork() system call



- Both the parent and the child execute **the same program**
- The child process starts its execution **at the location that `fork()` is returned**, not from the beginning of the program


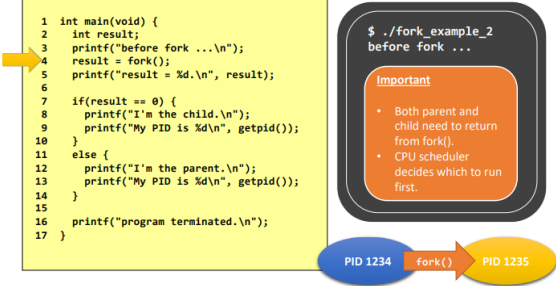
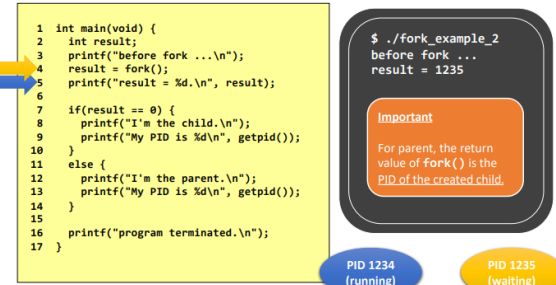
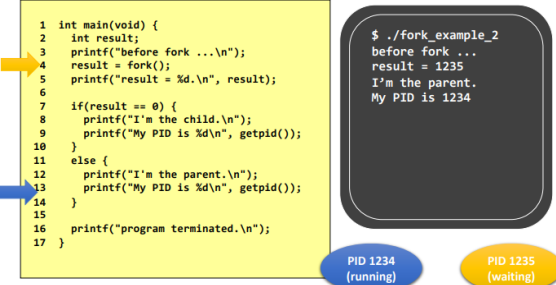
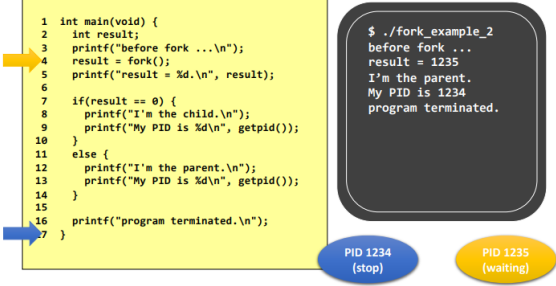
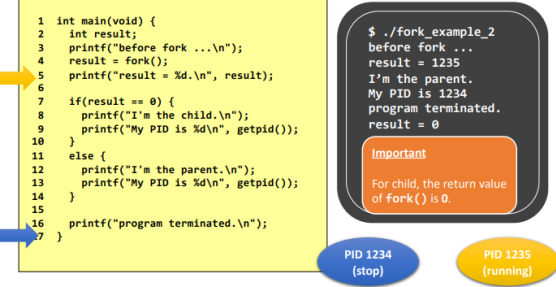
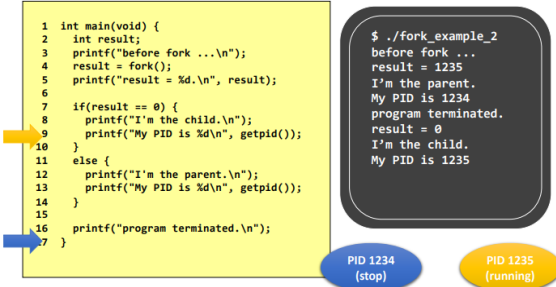
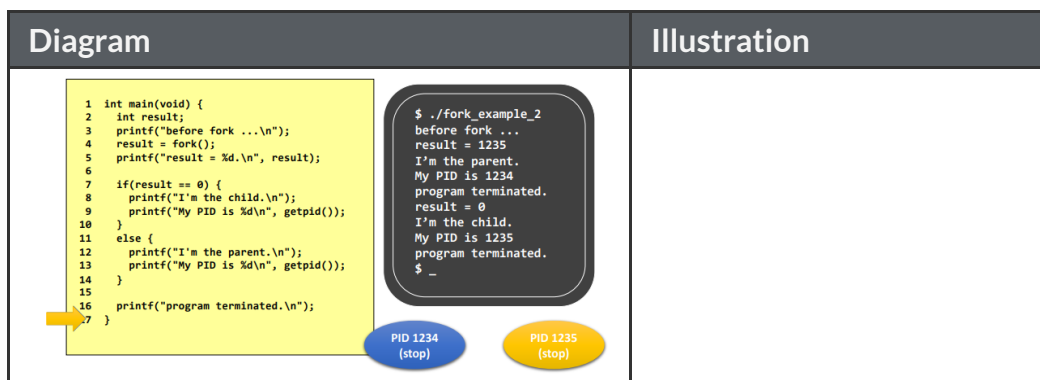
Diagram	Illustration
<pre>1 int main(void) { 2 int result; 3 printf("before fork ...\n"); 4 result = fork(); 5 printf("result = %d\n", result); 6 7 if(result == 0) { 8 printf("I'm the child.\n"); 9 printf("My PID is %d\n", getpid()); 10 } 11 else { 12 printf("I'm the parent.\n"); 13 printf("My PID is %d\n", getpid()); 14 } 15 16 printf("program terminated.\n"); 17 }</pre> 	

Diagram	Illustration
<pre> 1 int main(void) { 2 int result; 3 printf("before fork ...\n"); 4 result = fork(); 5 printf("result = %d.\n", result); 6 7 if(result == 0) { 8 printf("I'm the child.\n"); 9 printf("My PID is %d\n", getpid()); 10 } 11 else { 12 printf("I'm the parent.\n"); 13 printf("My PID is %d\n", getpid()); 14 } 15 16 printf("program terminated.\n"); 17 } </pre> 	<p>Both parent and child need to return from <code>fork()</code></p> <p>CPU scheduler decides which to run first</p>
<pre> 1 int main(void) { 2 int result; 3 printf("before fork ...\n"); 4 result = fork(); 5 printf("result = %d.\n", result); 6 7 if(result == 0) { 8 printf("I'm the child.\n"); 9 printf("My PID is %d\n", getpid()); 10 } 11 else { 12 printf("I'm the parent.\n"); 13 printf("My PID is %d\n", getpid()); 14 } 15 16 printf("program terminated.\n"); 17 } </pre> 	<p>For parent, the return value of <code>fork()</code> is the PID of the created child</p>
<pre> 1 int main(void) { 2 int result; 3 printf("before fork ...\n"); 4 result = fork(); 5 printf("result = %d.\n", result); 6 7 if(result == 0) { 8 printf("I'm the child.\n"); 9 printf("My PID is %d\n", getpid()); 10 } 11 else { 12 printf("I'm the parent.\n"); 13 printf("My PID is %d\n", getpid()); 14 } 15 16 printf("program terminated.\n"); 17 } </pre> 	
<pre> 1 int main(void) { 2 int result; 3 printf("before fork ...\n"); 4 result = fork(); 5 printf("result = %d.\n", result); 6 7 if(result == 0) { 8 printf("I'm the child.\n"); 9 printf("My PID is %d\n", getpid()); 10 } 11 else { 12 printf("I'm the parent.\n"); 13 printf("My PID is %d\n", getpid()); 14 } 15 16 printf("program terminated.\n"); 17 } </pre> 	
<pre> 1 int main(void) { 2 int result; 3 printf("before fork ...\n"); 4 result = fork(); 5 printf("result = %d.\n", result); 6 7 if(result == 0) { 8 printf("I'm the child.\n"); 9 printf("My PID is %d\n", getpid()); 10 } 11 else { 12 printf("I'm the parent.\n"); 13 printf("My PID is %d\n", getpid()); 14 } 15 16 printf("program terminated.\n"); 17 } </pre> 	<p>For child, the return value of <code>fork()</code> is 0</p>
<pre> 1 int main(void) { 2 int result; 3 printf("before fork ...\n"); 4 result = fork(); 5 printf("result = %d.\n", result); 6 7 if(result == 0) { 8 printf("I'm the child.\n"); 9 printf("My PID is %d\n", getpid()); 10 } 11 else { 12 printf("I'm the parent.\n"); 13 printf("My PID is %d\n", getpid()); 14 } 15 16 printf("program terminated.\n"); 17 } </pre> 	



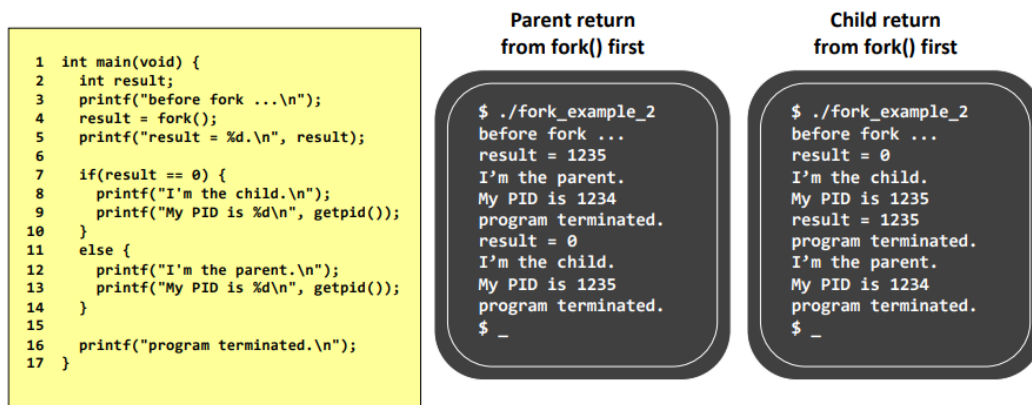
fork() system call

fork() does not clone the following

Distinct items	Parent	Child
Return value of fork()	PID of the child process.	0
PID	Unchanged.	Different, not necessarily be "Parent PID + 1"
Parent process	Unchanged.	Parent.
Running time	Cumulated.	Just created, so should be 0.
[Advanced] File locks	Unchanged.	None.

- fork() duplicates the current process itself
- If a process can only duplicate itself and always runs the same program, it's not quite meaningful

CPU Scheduler and fork()



exec() system call


```
int main(void) {
    printf("before execl ...\n");
    execl("/bin/ls", "/bin/ls", NULL);
    printf("after execl ...\n");
    return 0;
}
```

```
$/exec_example
before execl ...
exec_example
exec_example.c
```

What is the output?

The same as the output of running "ls" in the shell.

- `execl()`: a member of the exec system call family (and the family has 6 members)
 - 1st argument: program name `/bin/ls` in the example
 - 2nd argument: `argument[0]` to the program
 - 3rd argument: `argument[1]` to the program

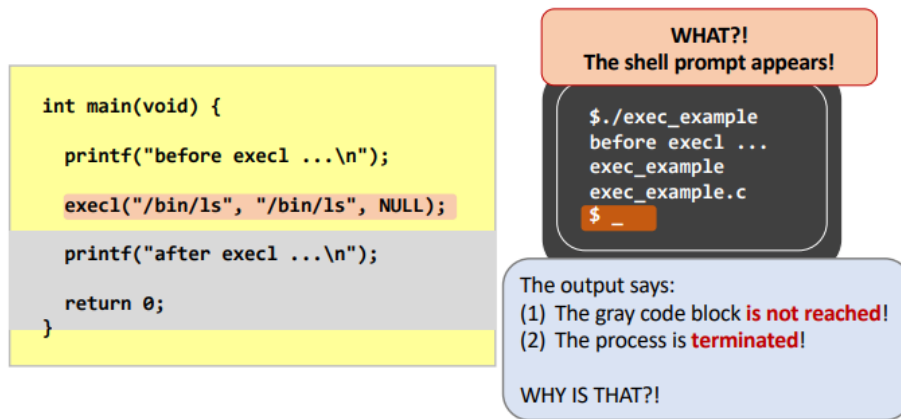
`execl("/bin/ls", "/bin/ls", NULL);`

Argument Order	Value in above example	Description
1	<code>"/bin/ls"</code>	The file that the programmer wants to execute.
2	<code>"/bin/ls"</code>	When the process switches to <code>"/bin/ls"</code> , this string is the program argument[0] .
3	<code>NULL</code>	This states the end of the program argument list.

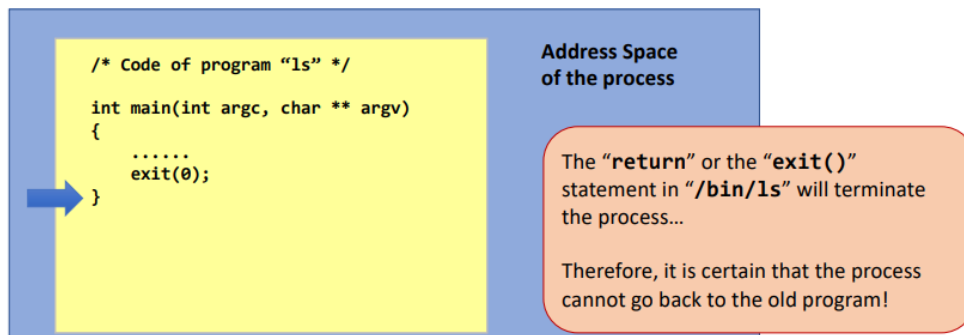
`execl("/bin/ls", "/bin/ls", "-l", NULL);`

Argument Order	Value in above example	Description
1	<code>"/bin/ls"</code>	The file that the programmer wants to execute.
2	<code>"/bin/ls"</code>	When the process switches to <code>"/bin/ls"</code> , this string is the program argument[0] .
3	<code>"-l"</code>	When the process switches to <code>"/bin/ls"</code> , this string is the program argument[1] .
4	<code>NULL</code>	This states the end of the program argument list.

- The `exec` system call family is not simply a function that “invokes” a command



- `exec()` loads program “ls” into the **memory** of this process



- The `return` or the `exit()` statement in `/bin/ls` will terminate the process
- Therefore, it is certain that the process **cannot go back to the old program**
- The process is changing the code that is executing and never returns to the original code
- The process that calls an `exec` system call will **replace user-space info**, e.g.,
 - **Program Code**
 - **Memory**: local variables, global variables, and dynamically allocated memory
 - **Register value**: e.g. PC
- But, the **kernel-space info** of that process is **preserved**, including
 - PID
 - Process relationship

wait() system call: Sync Parent with Child

```

1 int main(void) {
2     int result;
3     printf("before fork ...\n");
4     result = fork();
5     printf("result = %d.\n", result);
6
7     if(result == 0) {
8         printf("I'm the child.\n");
9         printf("My PID is %d\n", getpid());
10    }
11    else {
12        printf("I'm the parent.\n");
13        wait(NULL);
14        printf("My PID is %d\n", getpid());
15    }
16
17    printf("program terminated.\n");
18 }

```

Parent return
from fork() first

```

$ ./fork_example_2
before fork ...
result = 1235
I'm the parent.
result = 0
I'm the child.
My PID is 1235
program terminated.
My PID is 1234
program terminated.
$ _

```

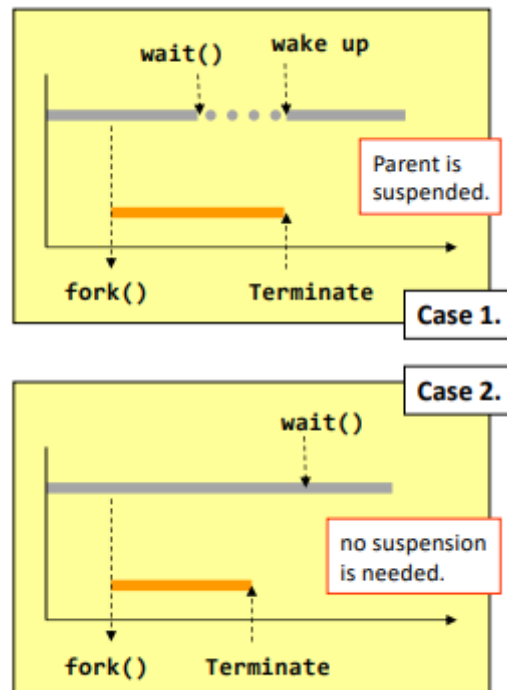
Child return
from fork() first

```

$ ./fork_example_2
before fork ...
result = 0
I'm the child.
My PID is 1235
result = 1235
program terminated.
I'm the parent.
My PID is 1234
program terminated.
$ _

```

- `wait()` suspends the calling process to waiting
- `wait()` returns when one of its child processes changes from running to terminated



- Return immediately (i.e., does nothing) if
 - It has **no children**
 - Or a child terminates before the parent calls `wait` for

`wait()` v.s. `waitpid()`

- `wait()`
 - Wait for any one of the child processes
 - Detect child termination only
- `waitpid()`
 - Depending on the parameters, `waitpid()` will wait for a **particular child** only

- Depending on the parameters, `waitpid()` can detect **different status changes of the child** (resume/stop by a signal)

3. Processes: Kernel view


Process Control Block (PCB)

process state
process number
program counter
registers
memory limits
list of open files
...

Information associated with each process

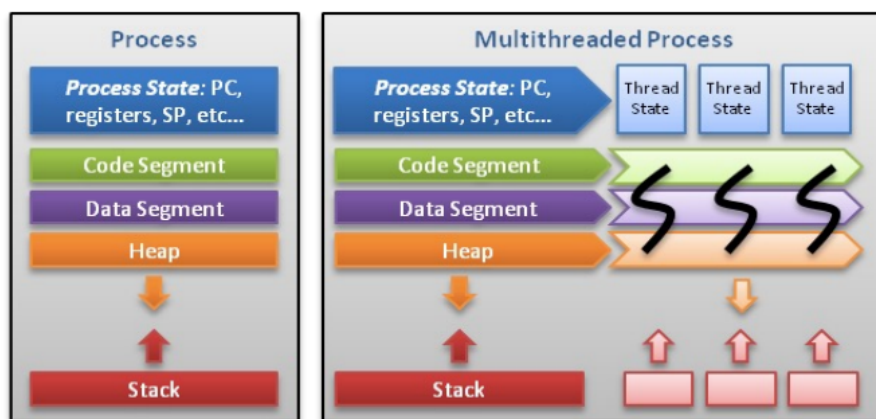
- Process state
- Program counter
- CPU registers
- CPU scheduling information
- Memory-management information
- Accounting information
- I/O status information

Ready Queue and I/O Device Queues



- PCBs are linked in multiple **queues**
 - **Ready queue** contains all processes in the *ready* state (to run on this CPU)
 - **Device queue** contains processes waiting for I/O events from this device
 - Process may migrate among these queues

Process and Threads



Threads contain only necessary information, such as a stack (for local variables, function arguments, return values), a copy of the registers, program counter and any thread-specific data to allow them to be scheduled individually. Other data is shared within the process between all threads.

© Alfred Park, <http://randu.org/tutorials/threads>

- One process may have **more than one threads**
 - A single-threaded process performs a single thread of execution
 - A multi-threaded process performs multiple threads of execution “concurrently”, thus allowing short response time to user’s input

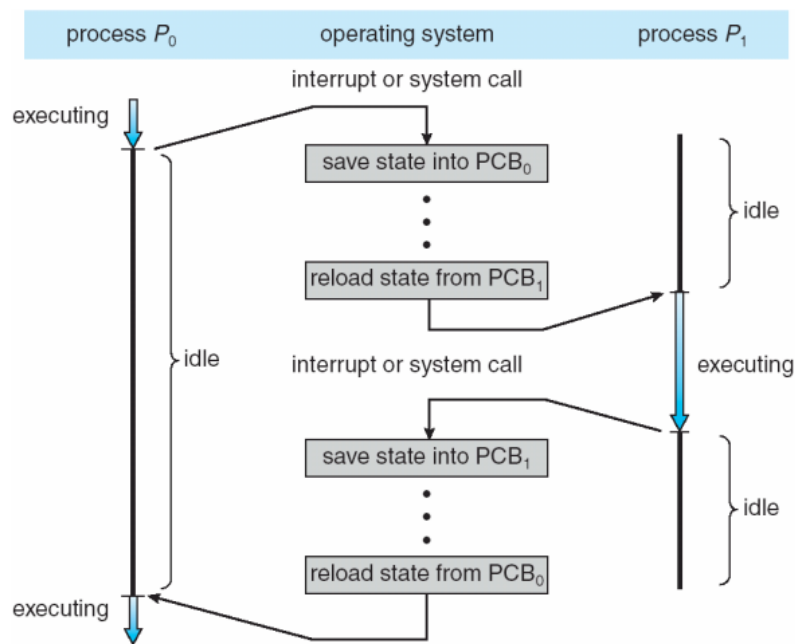
even when the main thread is busy

- PCB is extended to **include information about each thread**

Switching Between Processes

- Once a **process runs on a CPU**, it only gives back the control of a CPU
 - when it makes a **system call**
 - when it raises an **exception**
 - when an **interrupt** occurs
- What if none of these would happen for a long time?
 - Cooperative scheduling: OS will have to wait
 - Early Macintosh OS, old Alto system
 - Non-cooperative scheduling: timer interrupts
 - Modern operating systems
- When OS kernel regains the control of CPU
 - It first **completes** the task
 - Serve system call
 - Handle interrupt/exception
 - It then **decides which process to run next**
 - by asking its CPU scheduler
 - It performs a **context switch** if the soon-to-be-executing process is different from the previous one

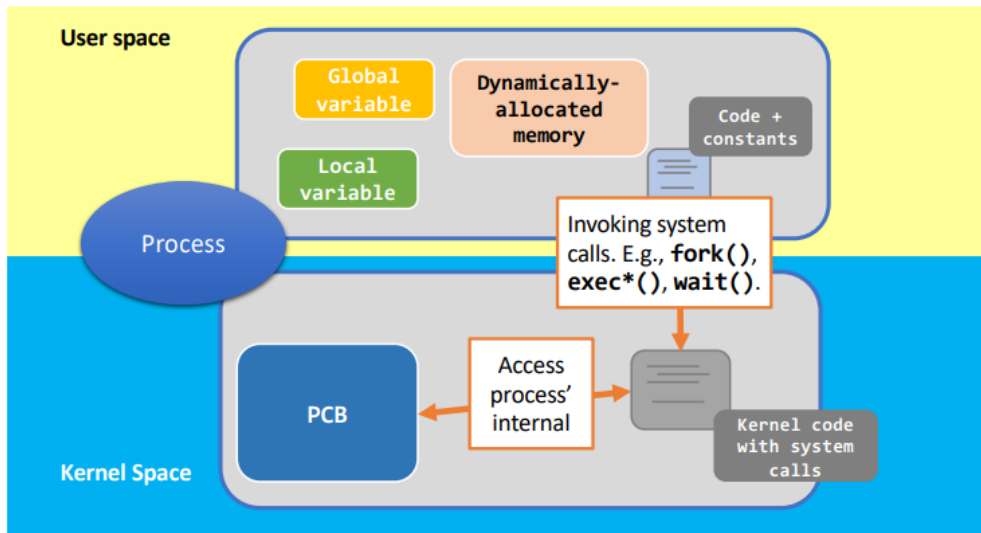
Context Switch



- During context switch, the system must **save the state of the old process** and **load the saved state for the new process**

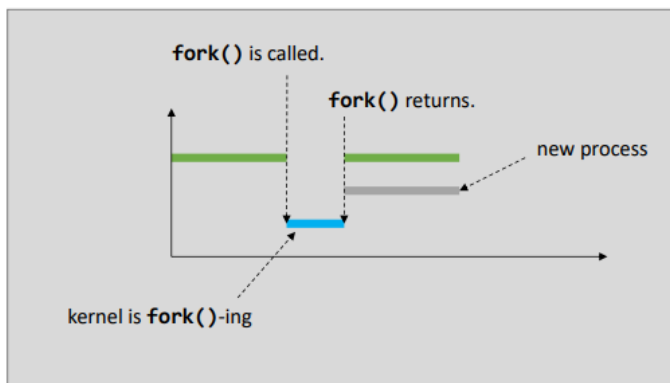
- Context of a process is represented in the **PCB**
- The time used to do context switch is an overhead of the system; the system does no useful work while switching
 - Time of context switch depends on **hardware** support
 - Context switch **cannot be too frequent**

4. Kernel view of fork(), exec() and wait()

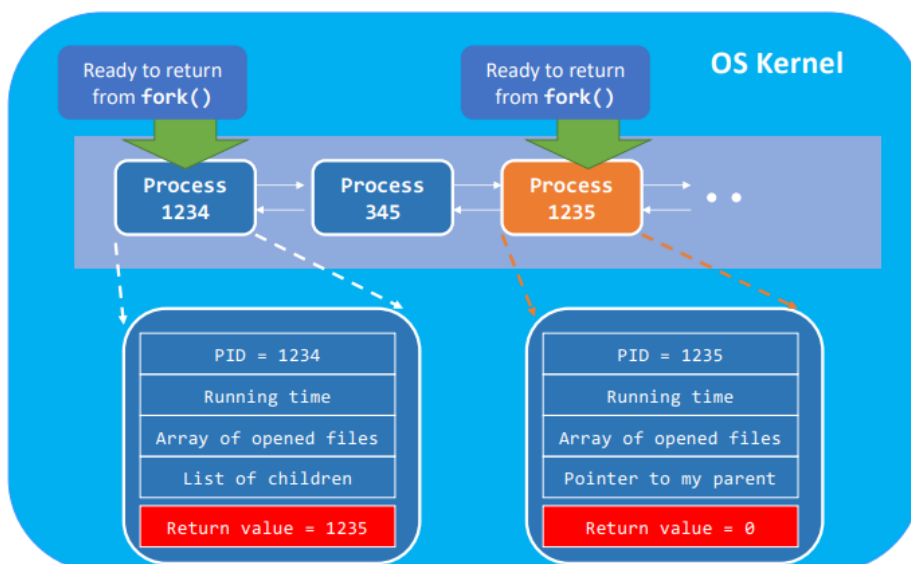
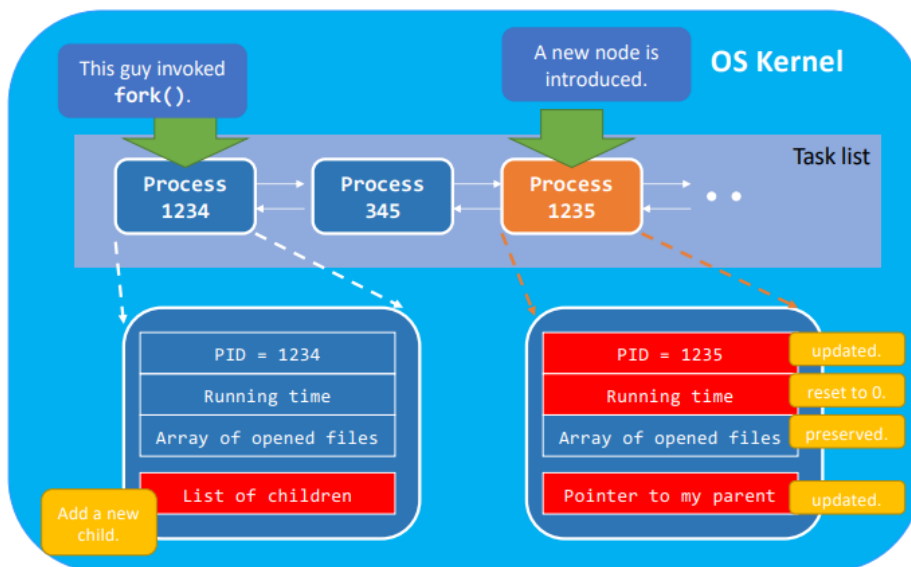
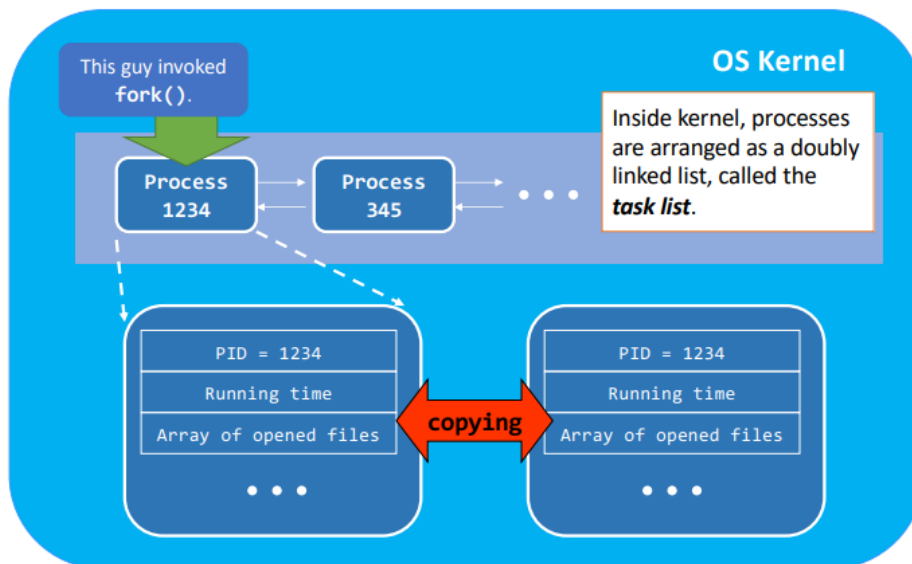


Fork

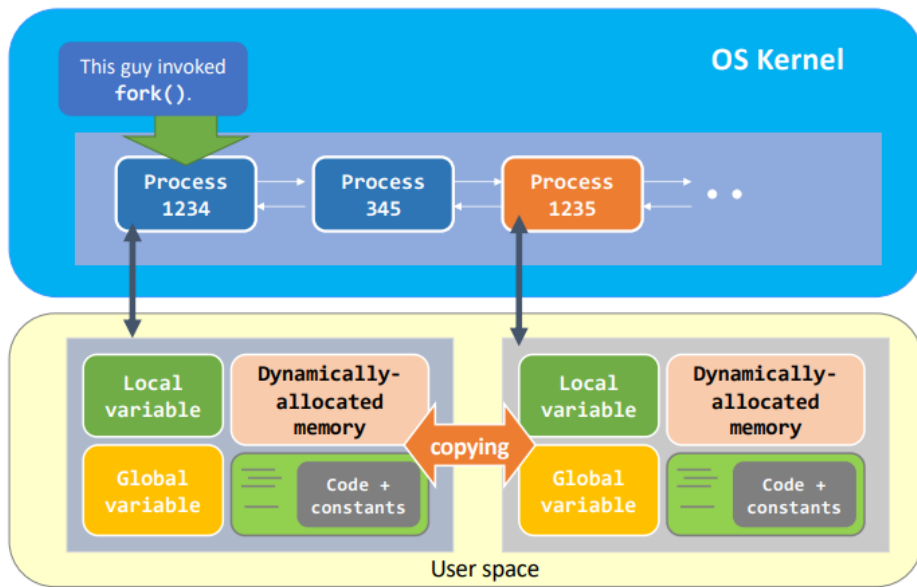
In User Mode



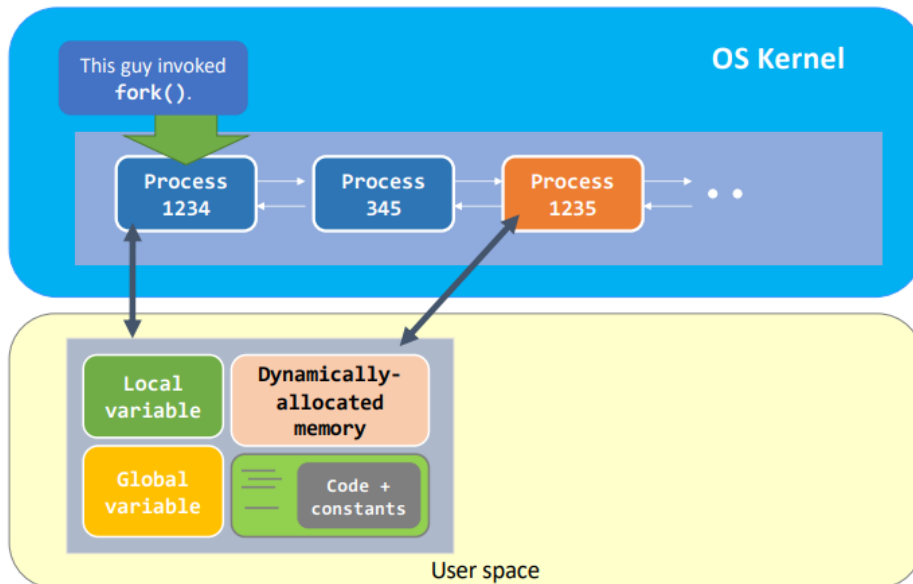
In Kernel Mode



Case 1: Duplicate Address Space

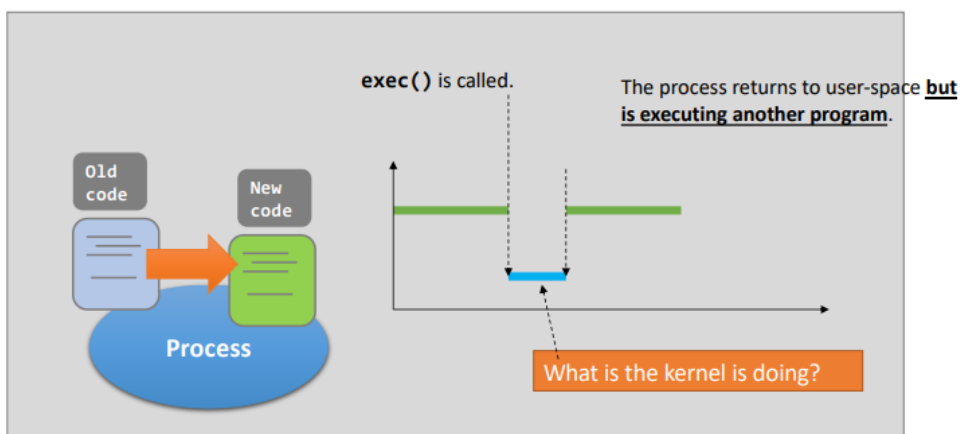


Case 2: Copy on Write

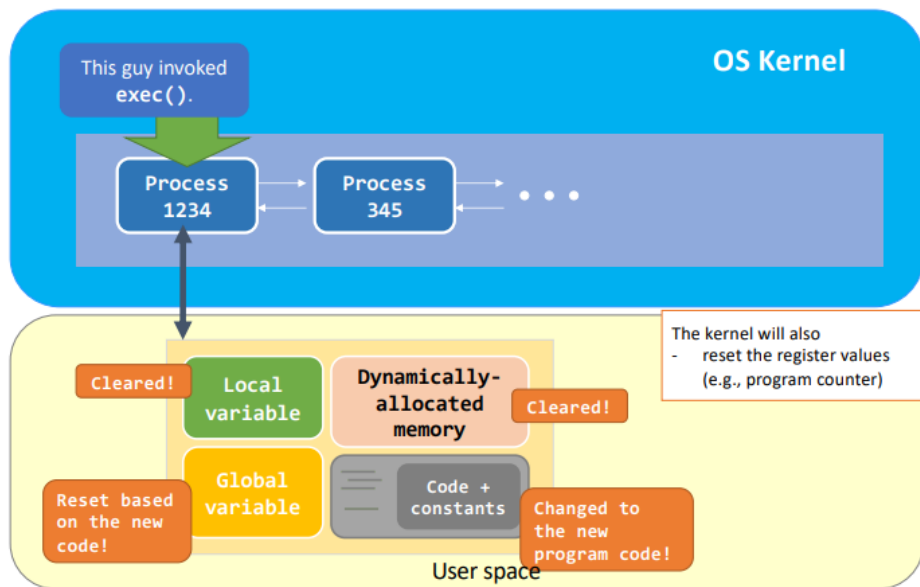


Exec

In User Mode

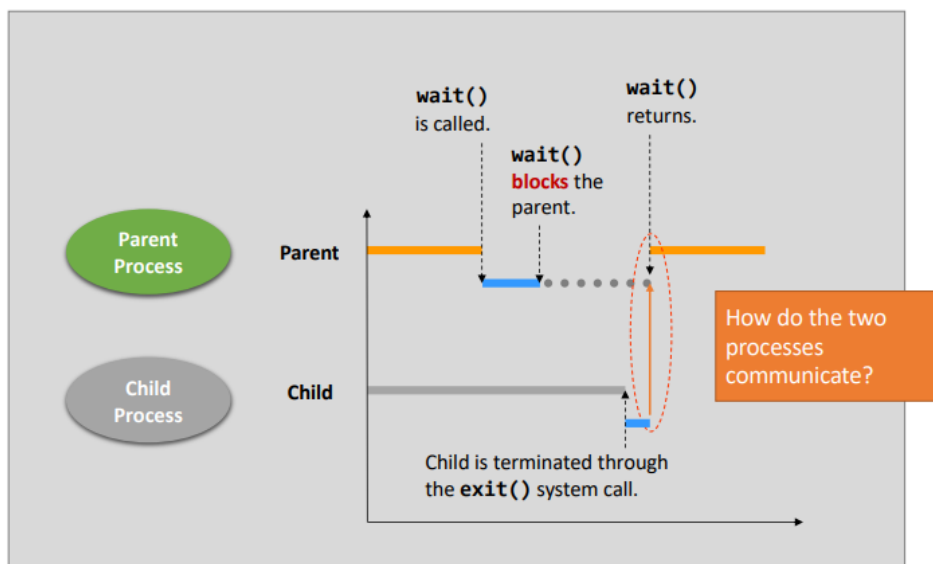


In Kernel Mode

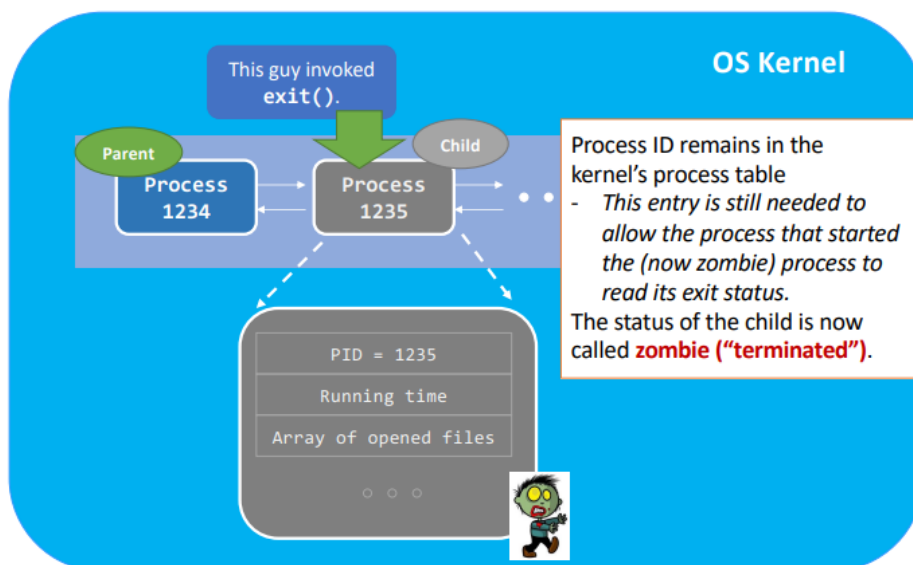
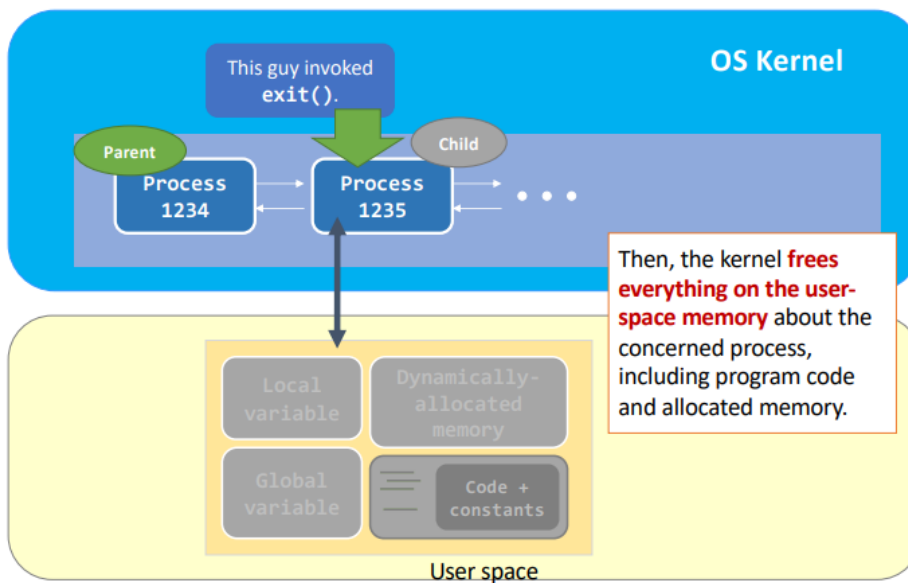
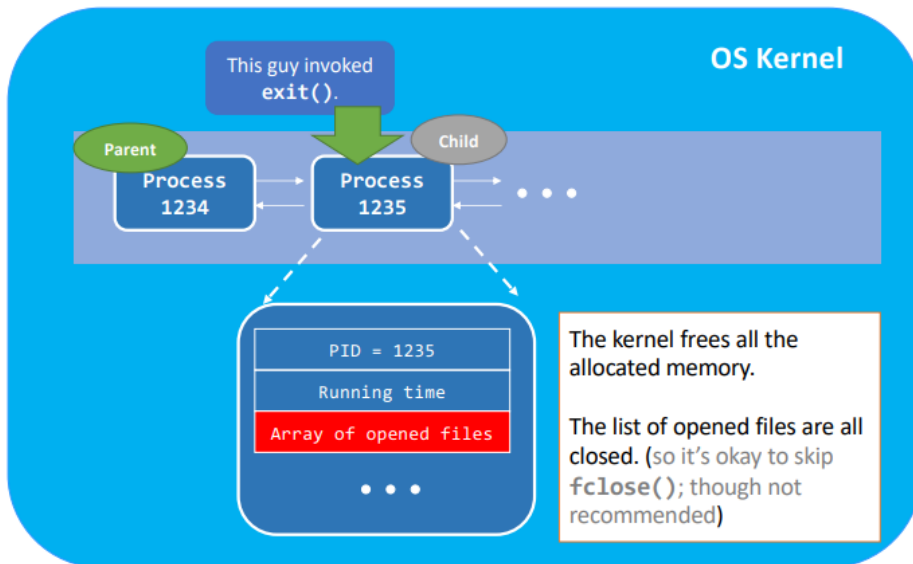


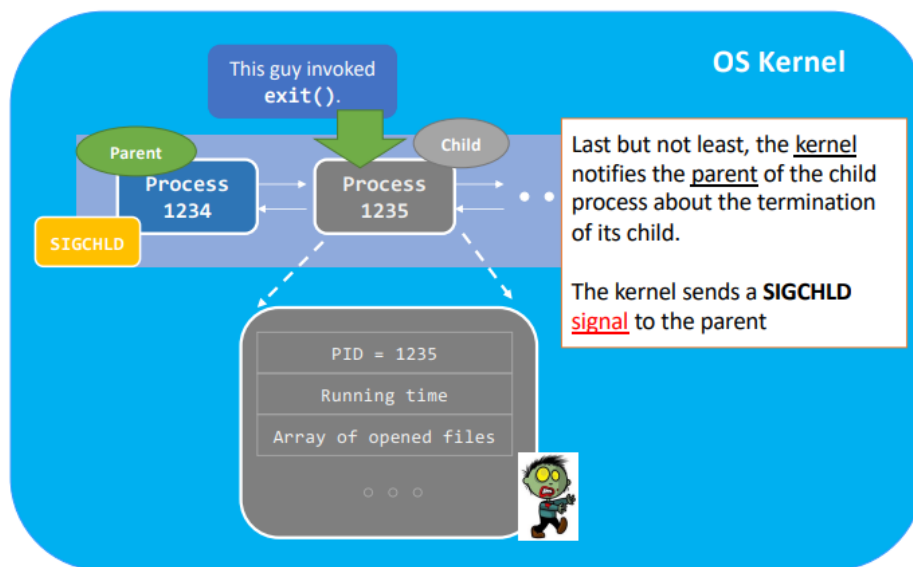
Exit

In User Mode



In Kernel Mode





Summary

Exit() will execute the following processes

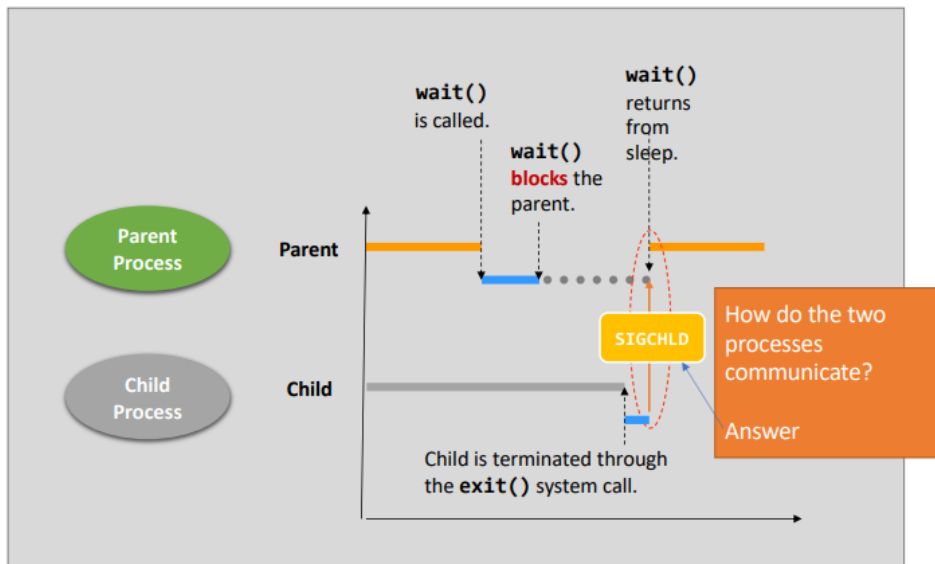
1. Clean up **most of** the allocated **kernel-space memory** (e.g., process's running time info).
2. Clean up the exit process's **user-space memory**
3. Notify the parent with `SIGCHLD`

Exit() system call turns a process into a zombie when

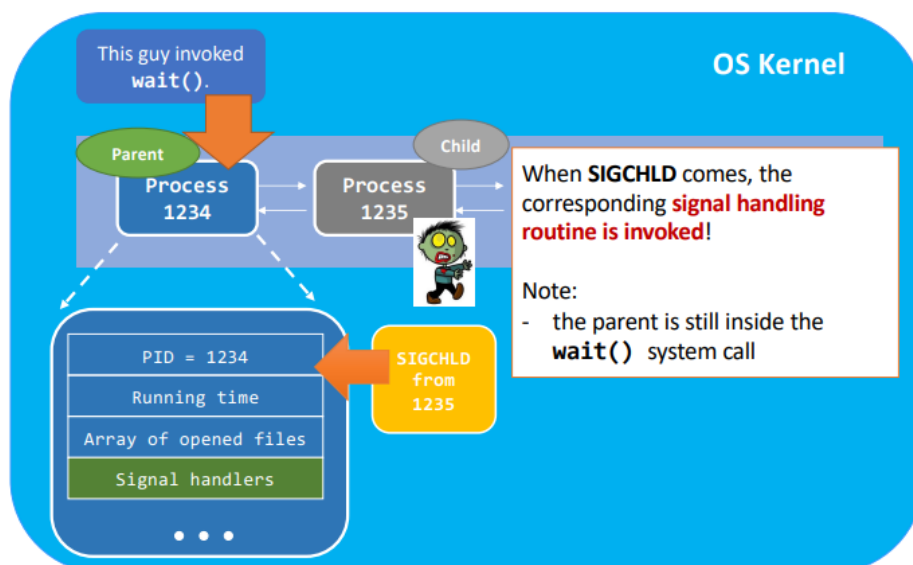
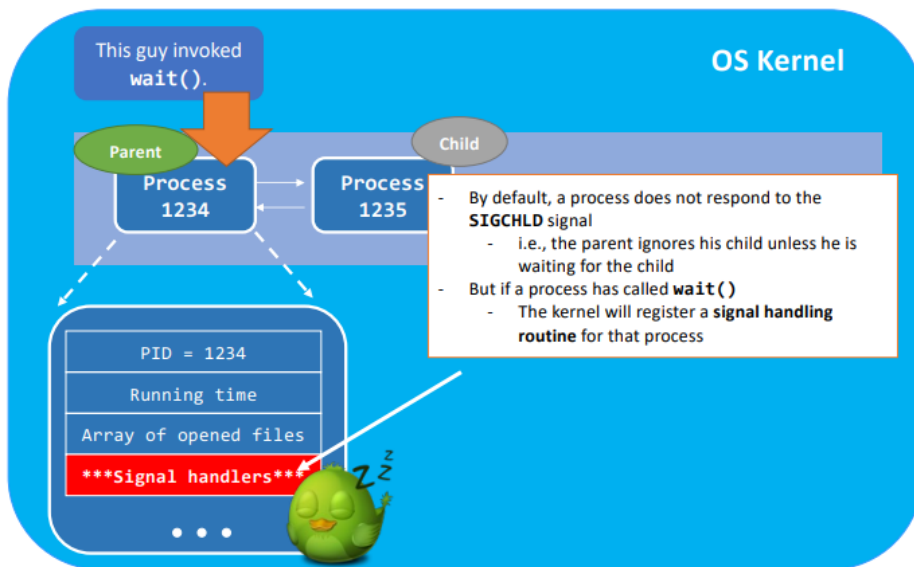
- The process calls `exit()`
- The process returns from `main()`
- The process terminates abnormally
 - The kernel knows that the process is terminated abnormally. Hence, the kernel invokes `exit()` for it

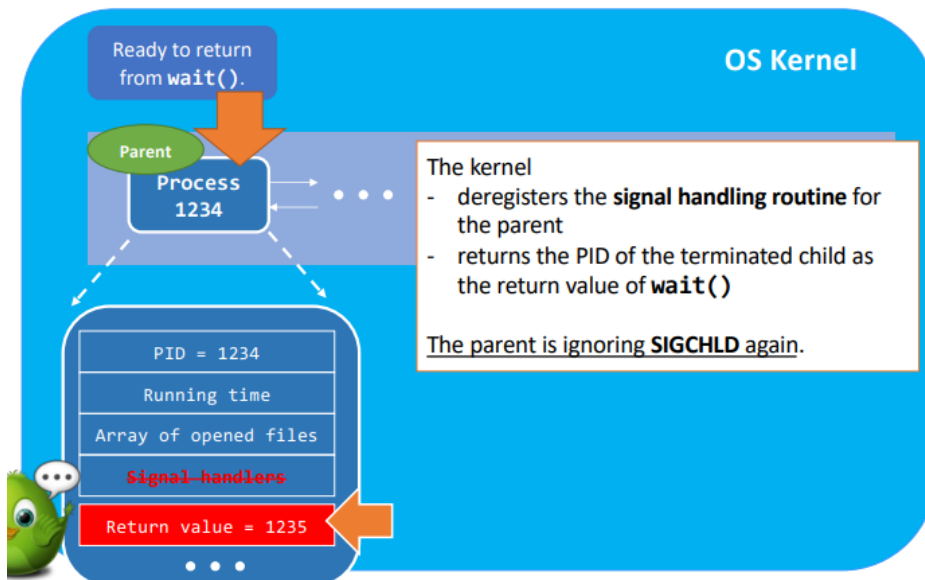
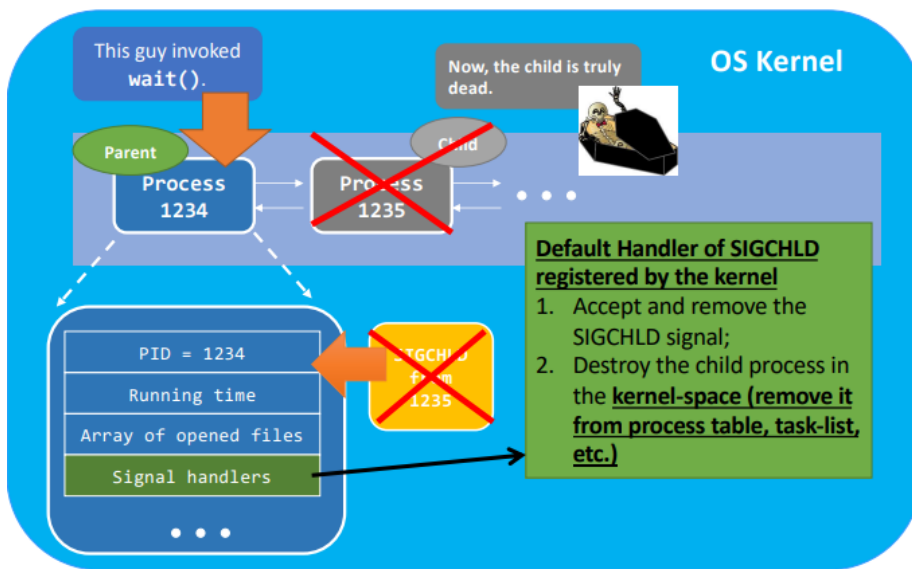
Wait

In User Mode

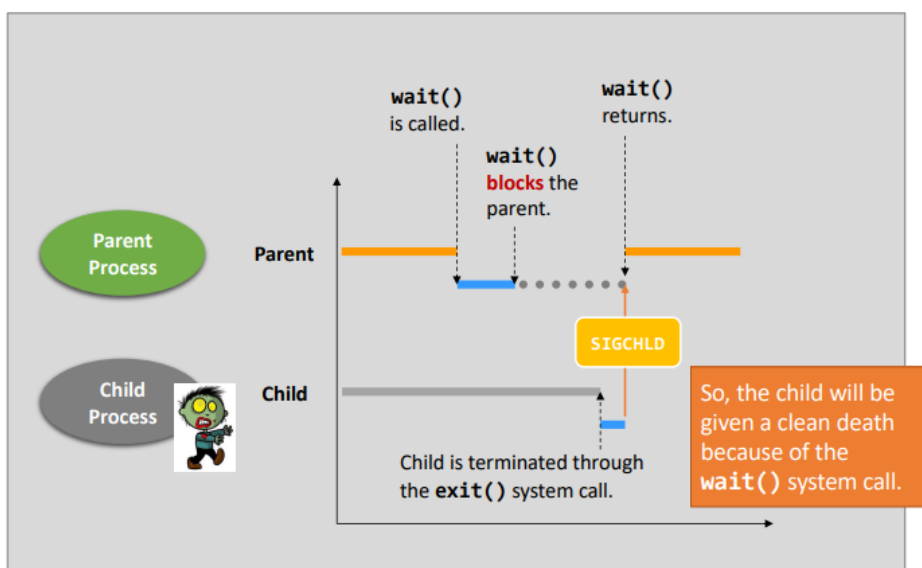


In Kernel View



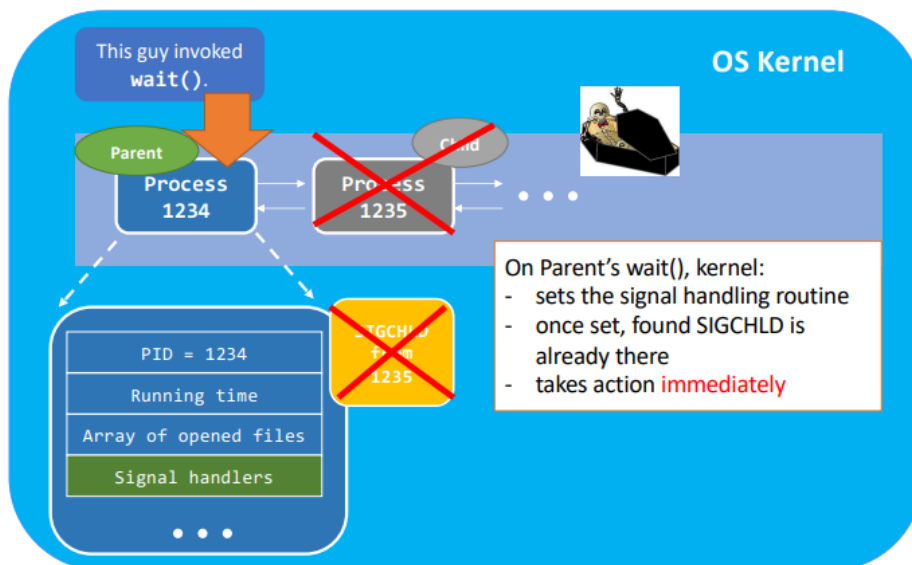
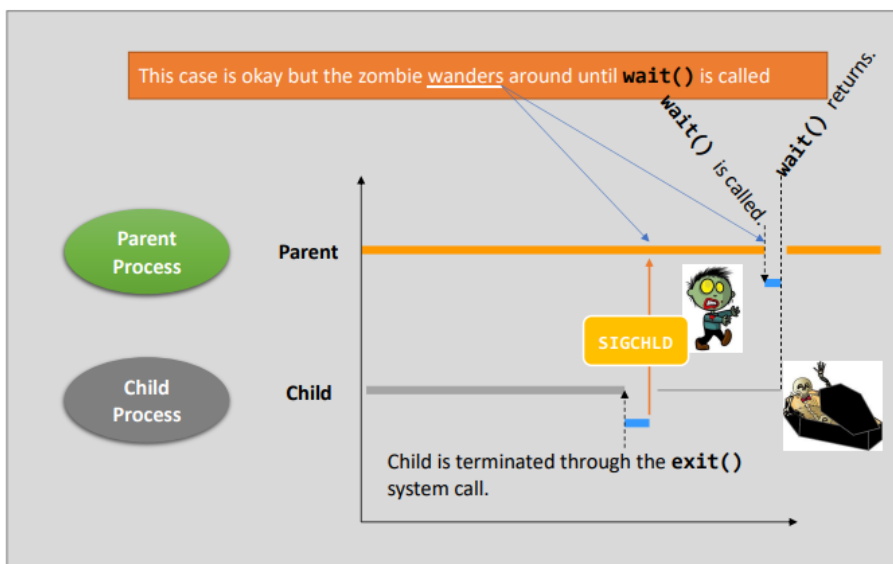
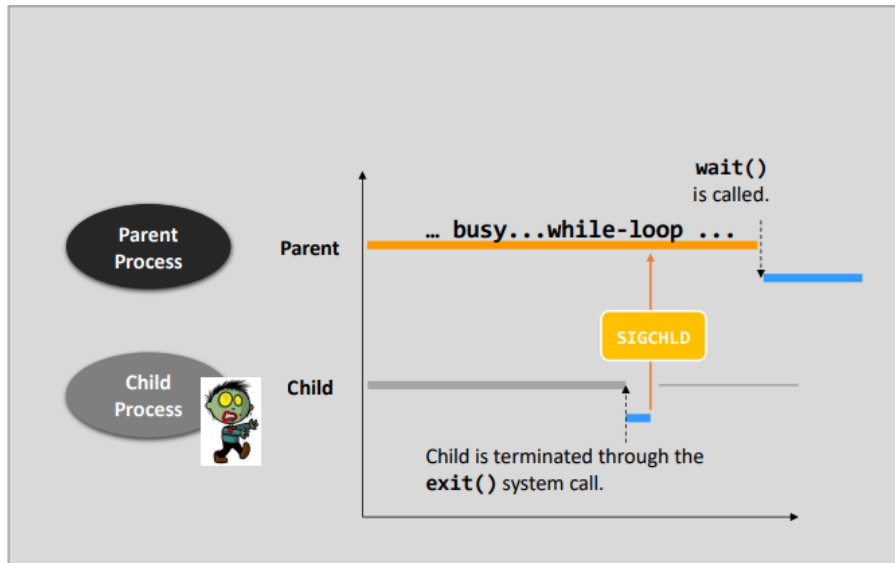


Normal Case



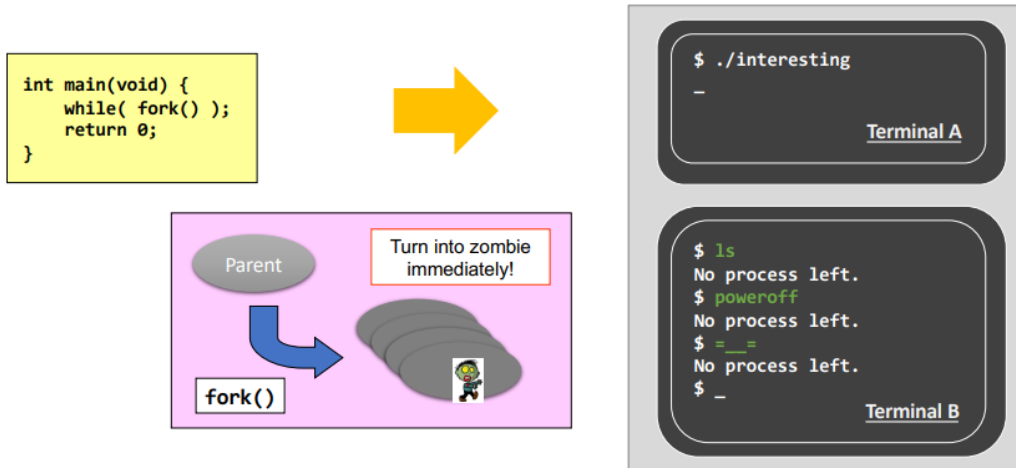
Abnormal Case

Parent's wait() after Child's exit()



Using wait() for Resource Management

- It is not only about process execution / suspension
- It is about system resource management
 - A zombie takes up a PID
 - The total number of PIDs are limited



Summary

- `wait()` & `waitpid()` reap zombie child processes
 - It is a must that you should never leave any zombies in the system
 - `wait()` & `waitpid()` pause the caller until
 - A child terminates/stops
 - The caller receives a signal (i.e., the signal interrupted the `wait()`)
- Linux will label zombie processes as `<defunct>`

```
$ ps aux | grep defunct
..... 3150 ... [ls] <defunct>
$ _
```

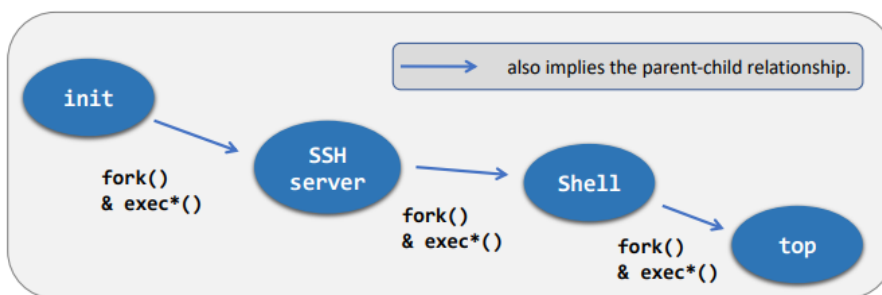
PID of the process

5. More about processes

The first process

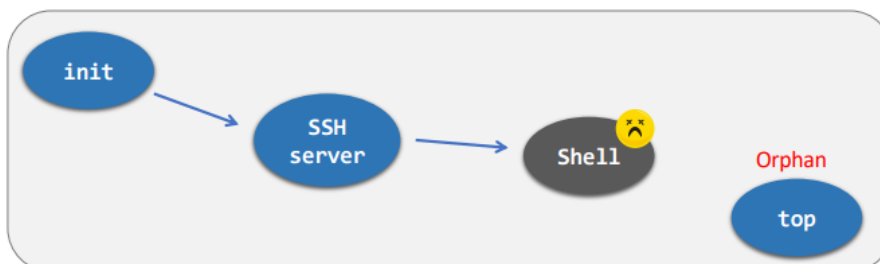
- We now focus on the process-related events
 - The kernel, while it is booting up, creates **the first process** - `init`
- The `init` process
 - has PID=1
 - is running the program code `/sbin/init`
- Its first task is to create more processes
 - using `fork()` and `exec()`

A Tree of Processes



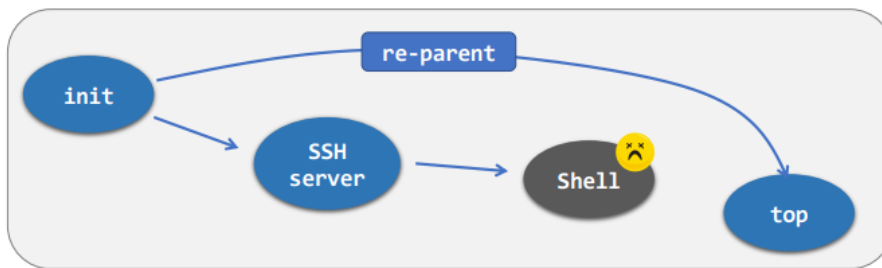
- You can view the tree with the command:
 - `ps tree`
 - `ps tree -A`: for ASCII-character-only display

Orphans



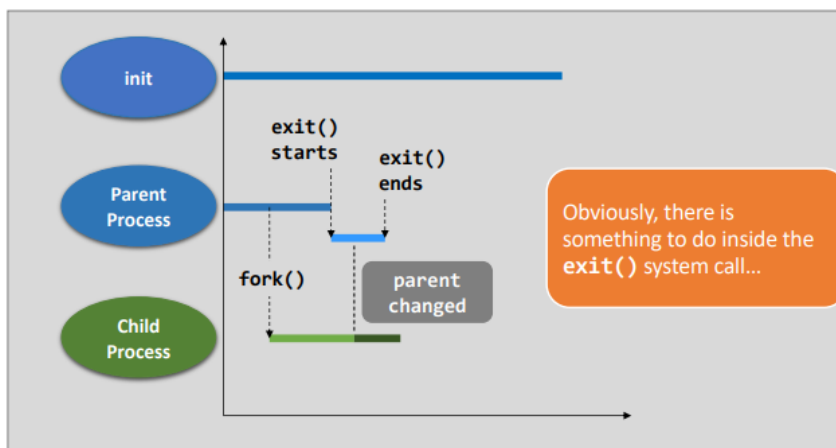
- However, termination can happen, at any time and in any place
 - This is no good because an orphan turns the hierarchy from a tree into a forest
 - no one would know the termination of the orphan

Re-parent

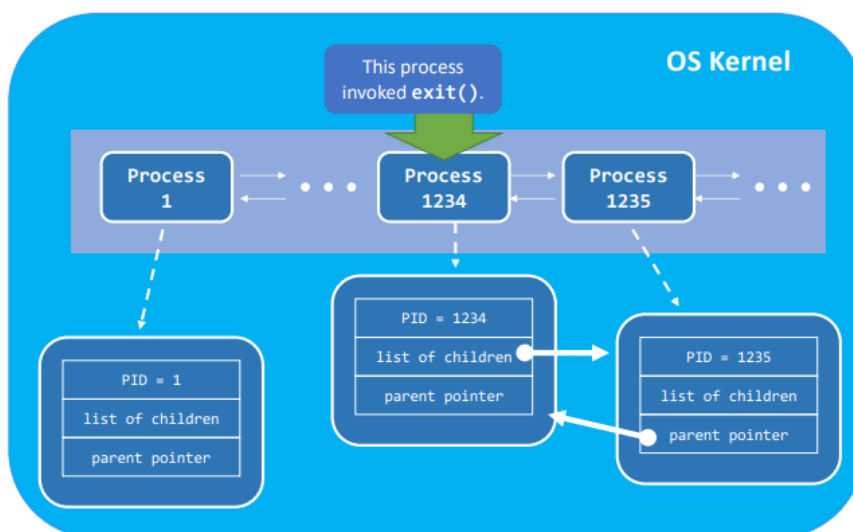


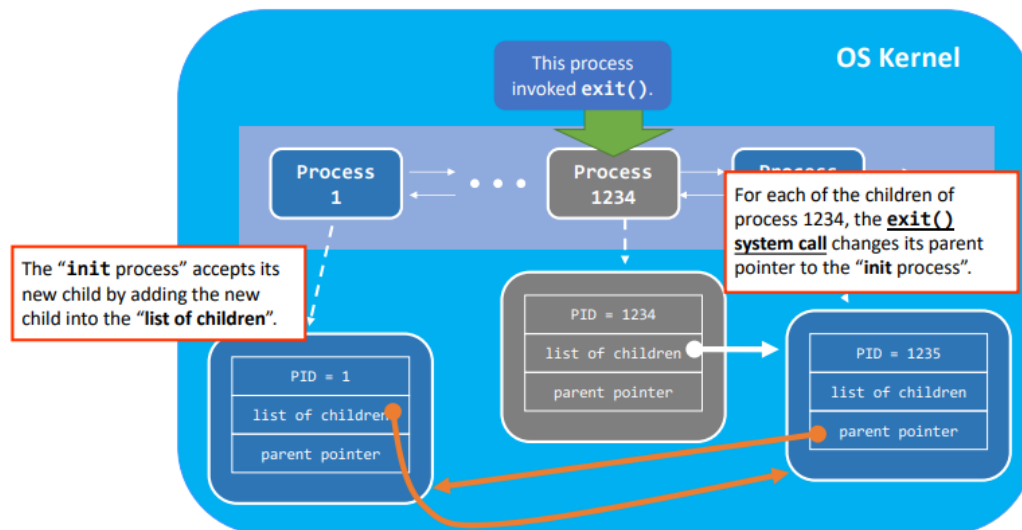
- In Linux
 - The `init` process will become the step-mother of all orphans
 - It's called **re-parenting**
- In Windows
 - It maintains a forest-like process hierarchy

In User Mode



In Kernel Mode

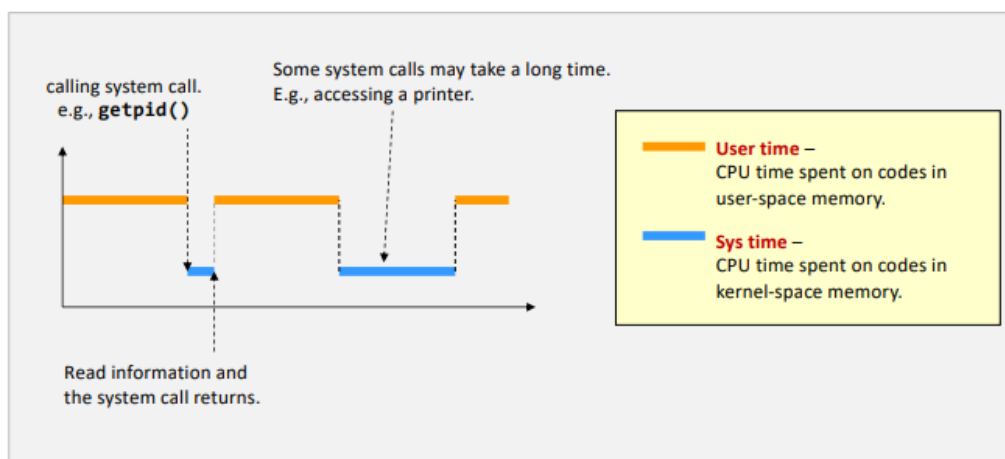




Background Jobs

- The re-parenting operation enables something called **background jobs** in Linux
 - It allows a process runs **without a parent terminal/shell**

Measure Process Time



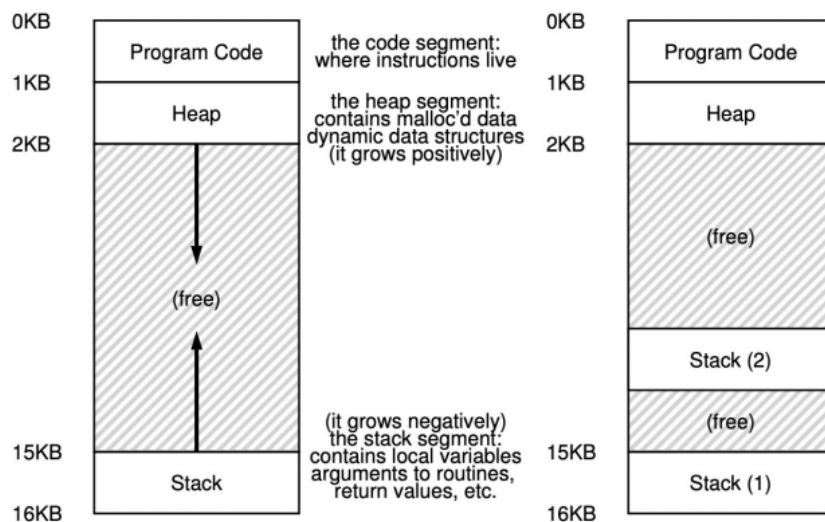
- The user time and the sys time together define **the performance of an application**
 - When writing a program, you must consider both the user time and the sys time
- Ideally, Real time = User time + Sys time
 - $\text{real} > \text{user} + \text{sys}$: **I/O intensive**
 - $\text{real} < \text{user} + \text{sys}$: **multi-core**

6. Threads

What is a Thread

- Thread is an **abstraction** of the **execution of a program**
 - A single-threaded program has one point of execution
 - A multi-threaded program has more than one points of execution
- Each thread has its own **private** execution state
 - Program counter and a private set of register
 - A private stack for thread-local storage
 - CPU switching from one thread to another requires context switch
- Threads in the same process **share** computing resources
 - Address space, files, signals, etc.

Single-Threaded and Multi-Threaded



Why Use Thread


- Increase **parallelism**
 - One thread per CPU makes better use of multiple CPUs to improve efficiency
- Avoid **blocking problem** progress due to slow I/O
 - Threading enables overlap of I/O with other activities within a single program
 - many modern server-based applications (web servers, database management systems, and the like) make use of threads
- Allow **resource sharing**

Thread Implementation


- **User-level thread**
 - Thread management (e.g., creating, scheduling, termination) done by user-level threads library

- OS does not know about user-level thread
- **Kernel-level thread**
 - Thread management done by kernel
 - OS is aware of each kernel-level thread


Thread Models



One-to-One



Many-to-One



Many-to-Many

- **One-to-one mapping**
 - One user-level thread to one kernel-level thread
 - Pros: Every thread can run or block independently
 - Cons: Need to make a crossing into kernel mode to schedule
- **Many-to-one mapping**
 - Many user-level thread to one kernel-level thread
 - Pros: context switch between threads is cheap
 - Cons: When one thread blocks on I/O, all threads block
- **Many-to-many mapping**
 - Many user-level thread to many kernel-level thread
 - Pros: best of the two worlds, more flexible
 - Cons: difficult to implement