

Lecture4 CPU Scheduling

1. CPU Scheduling

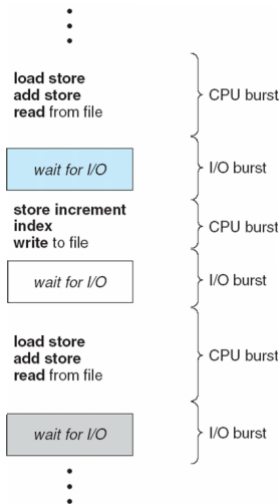

- Scheduling is important when multiple processes wish to run on a single CPU
 - CPU scheduler decides which process to run next

Two types of processes

	CPU-bound Process	I/O-bound process
Characteristic	Spends most of its running time on the CPU user-time > sys-time	Spends most of its running time on I/O sys-time > user-time
Example	AI course assignments	/bin/lis, networking programs

CPU Burst

- Process execution consists of a cycle of CPU execution and I/O wait
- CPU burst distribution



CPU Scheduler

- CPU scheduler selects one of the processes that are ready to execute and allocates the CPU to it
- CPU scheduling decisions may take place when a process
 1. Switches from **running** to **waiting** state
 2. Switches from **running** to **ready** state
 3. Switches from **waiting** to **ready**
 4. Terminates
- A scheduling algorithm takes place **only** under circumstances 1 and 4 is **non-preemptive**
- All other scheduling algorithms are **preemptive**

2. Scheduling Algorithm Optimization Criteria

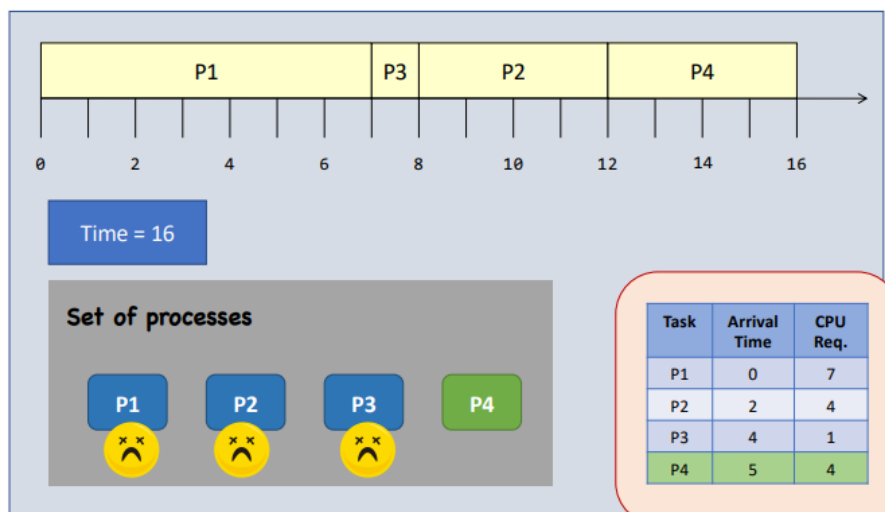
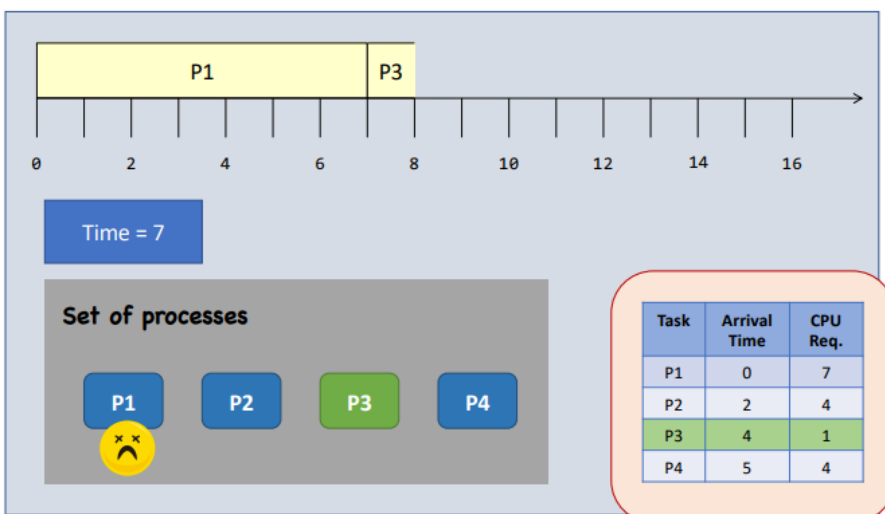
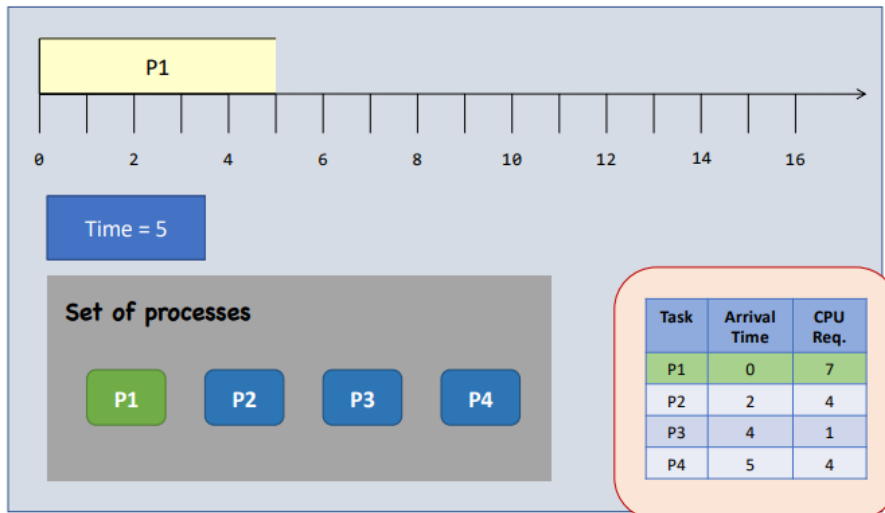
- Given a set of processes, with
 - **Arrival time:** the time they arrive in the CPU ready queue (from waiting state or from new state)
 - **CPU requirement:** their expected CPU burst time
- Minimize average **turnaround** time: The time between the arrival of the task and the time it is blocked or terminated

- Minimize average **waiting** time: The accumulated time that a task has waited in the ready queue
- Reduce the number of **context switches**

Shortest-job-first (SJF)

Non-preemptive SJF

Non-preemptive SJF will execute one process until it **finishes its CPU request**

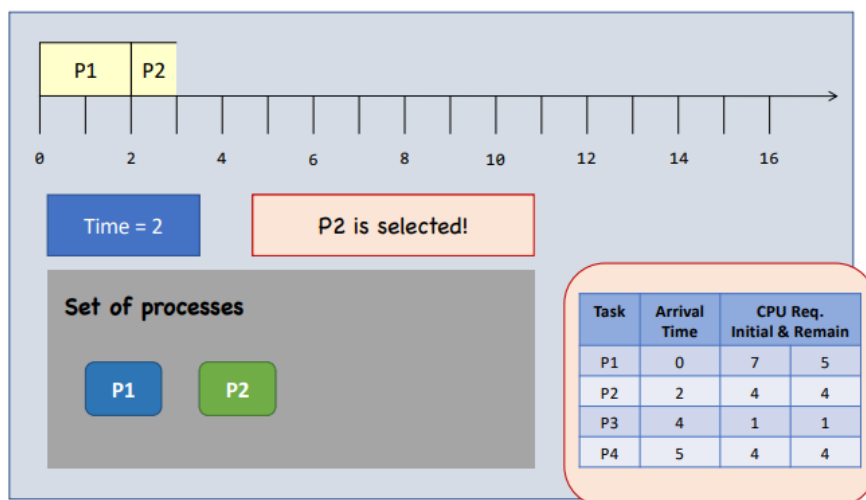



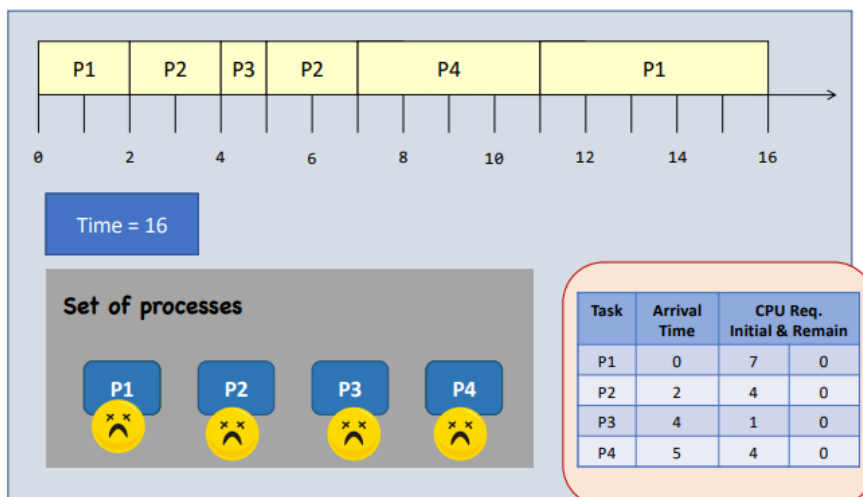
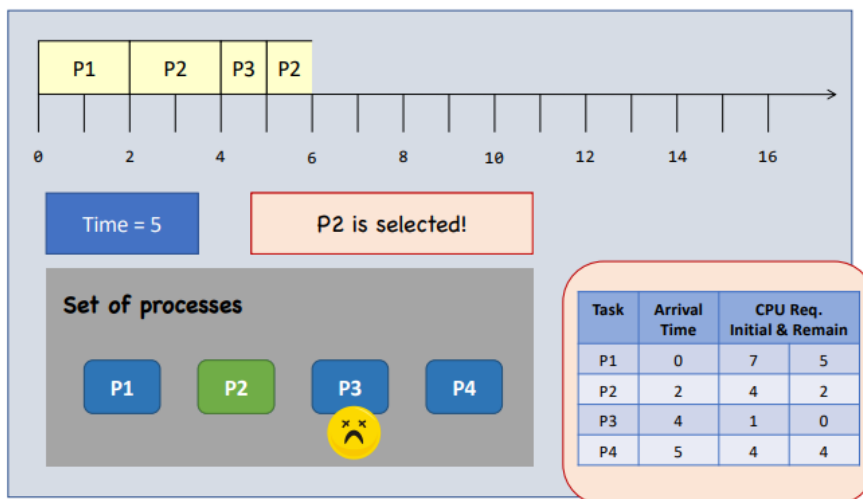
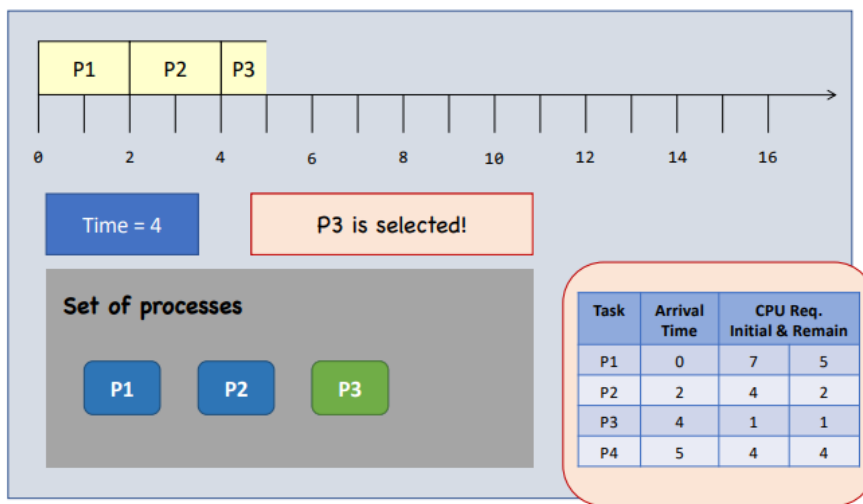
- waiting time:
 - $P1 = 0, P2 = 6, P3 = 3, P4 = 7$
 - $Average = (0 + 6 + 3 + 7) / 4 = 4$
- turnaround time:
 - $P1 = 7, P2 = 10, P3 = 4, P4 = 11$
 - $Average = (7 + 10 + 4 + 11) / 4 = 8$

Preemptive SJF

Whenever a new process arrives in the ready queue (either from waiting or from new state), the scheduler steps in and selects the next task based on **their remaining CPU requirements**

Suppose the time interrupt = 2





- waiting time
 - $P1 = 9, P2 = 1, P3 = 0, P4 = 2$
 - $\text{Average} = (9 + 1 + 0 + 2) / 4 = 3$
- turnaround time
 - $P1 = 16, P2 = 5, P3 = 1, P4 = 6$
 - $\text{Average} = (16 + 5 + 1 + 6) / 4 = 7$

SJF: Preemptive or Not?

	Non-preemptive SJF	Preemptive SJF
Average waiting time	4	3 (smallest)
Average turnaround time	8	7 (smallest)
# of context switching	3	5 (largest)

The waiting time and the turnaround time decrease at the expense of the increased number of context switches.

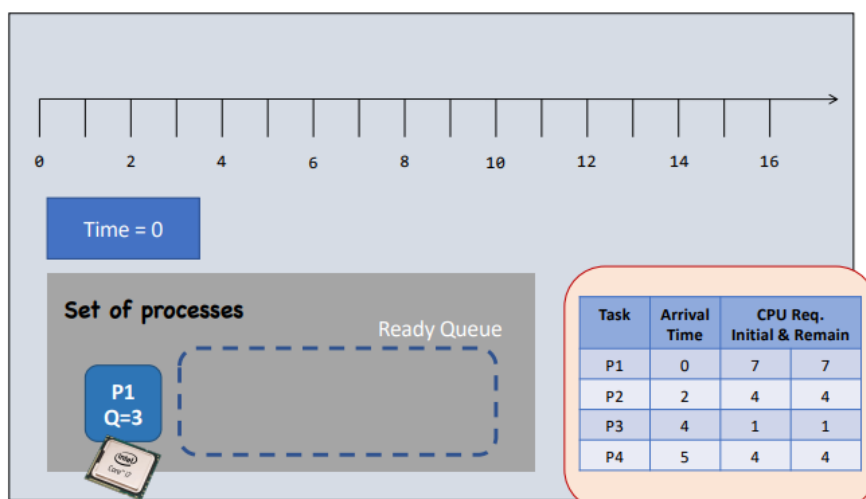
Task	Arrival Time	CPU Req.
P1	0	7
P2	2	4
P3	4	1
P4	5	4

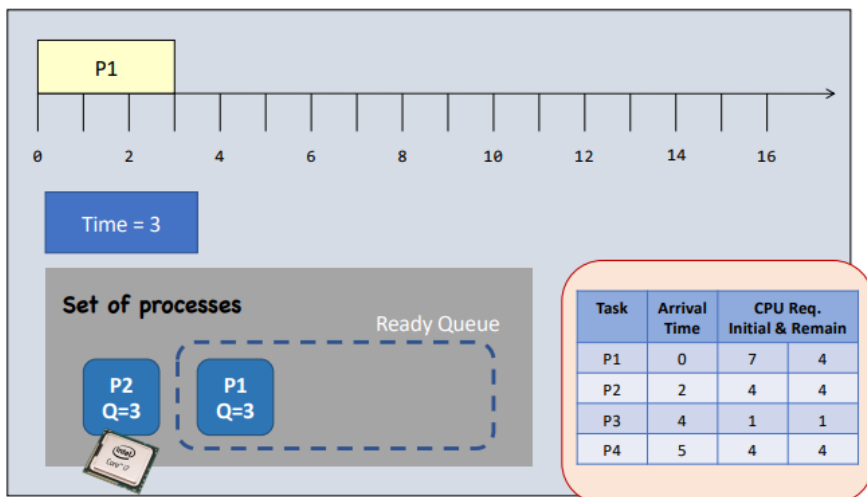
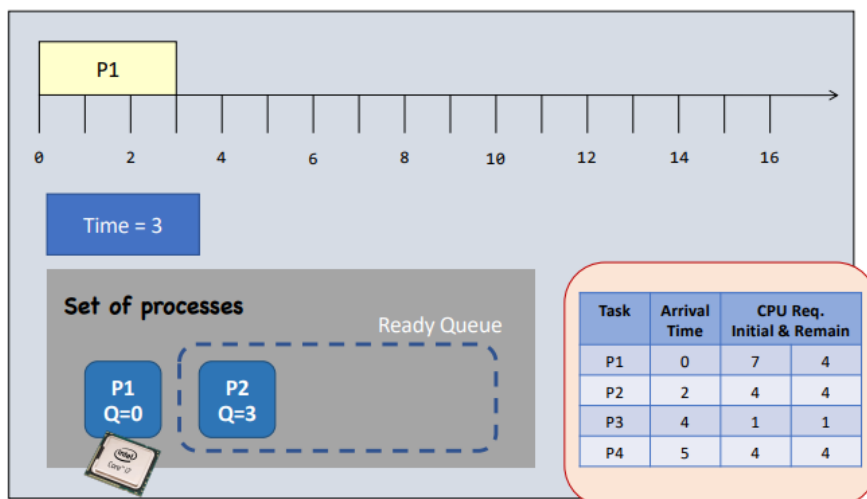
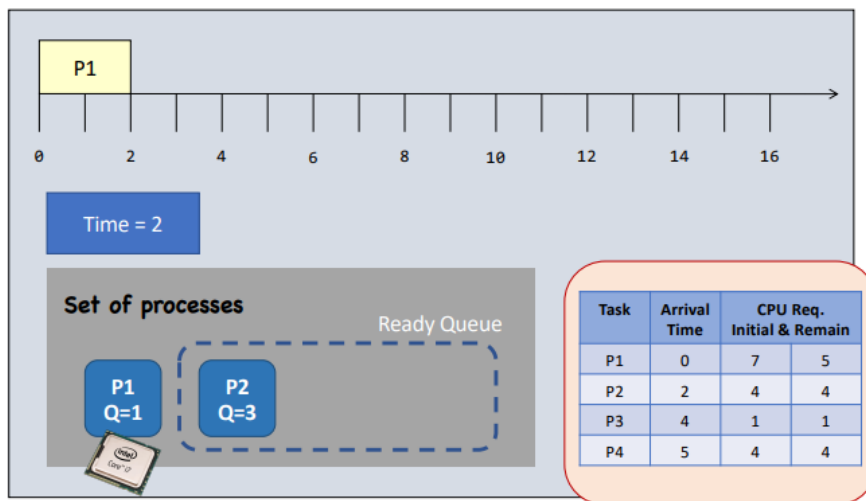
Round-robin (RR)

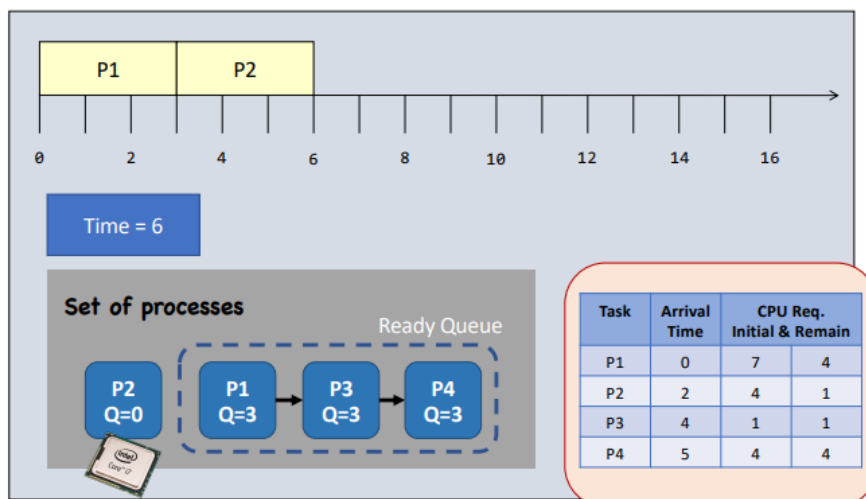
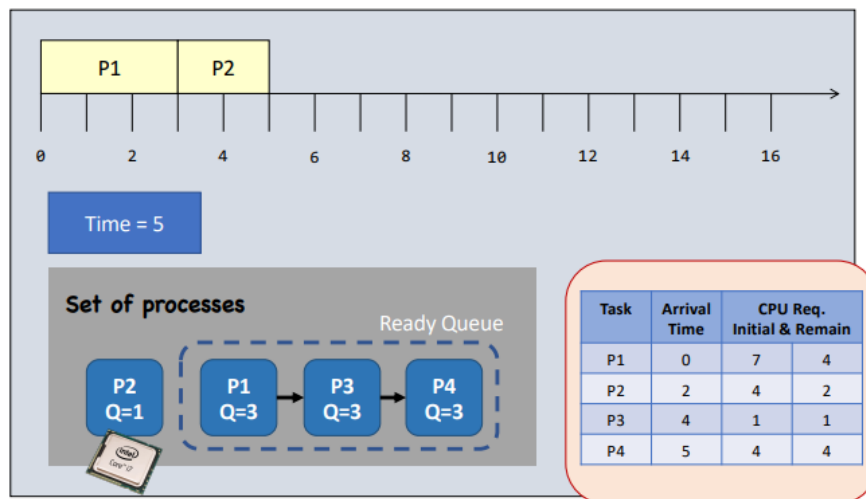
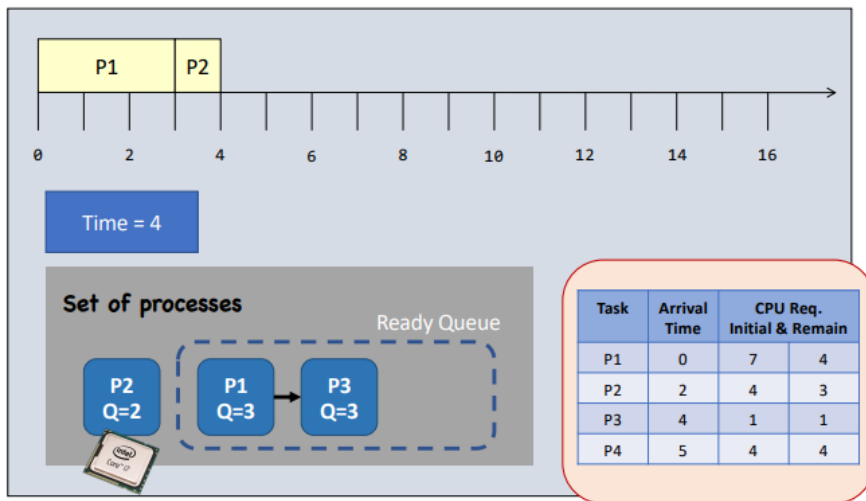
- Round-Robin (RR) scheduling is preemptive
 - Every process is given a **quantum** (the amount of time allowed to execute)
 - Whenever the quantum of a process is used up (i.e., 0), the process is preempted, placed at the end of the queue, with its quantum recharged
 - Then, the scheduler steps in and it chooses the next process which has a non-zero quantum to run
 - Processes are therefore running one-by-one as a circular queue
- New processes are **added to the tail** of the ready queue
 - New process's arrival won't trigger a new selection decision

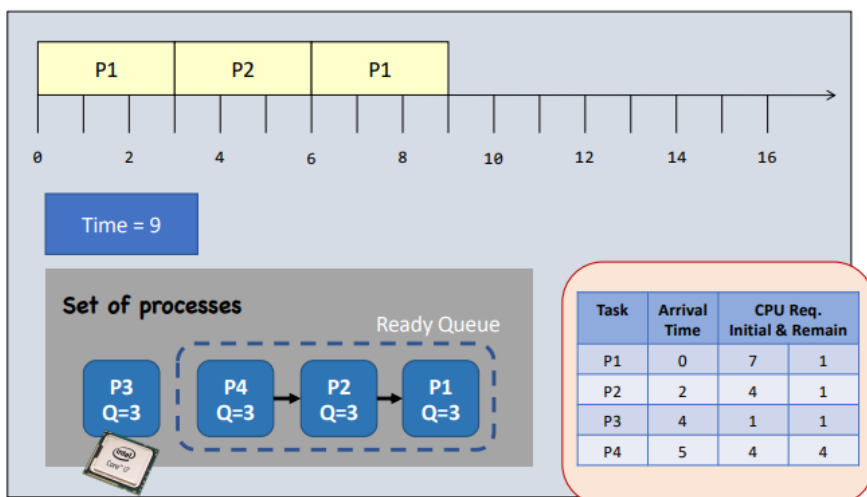
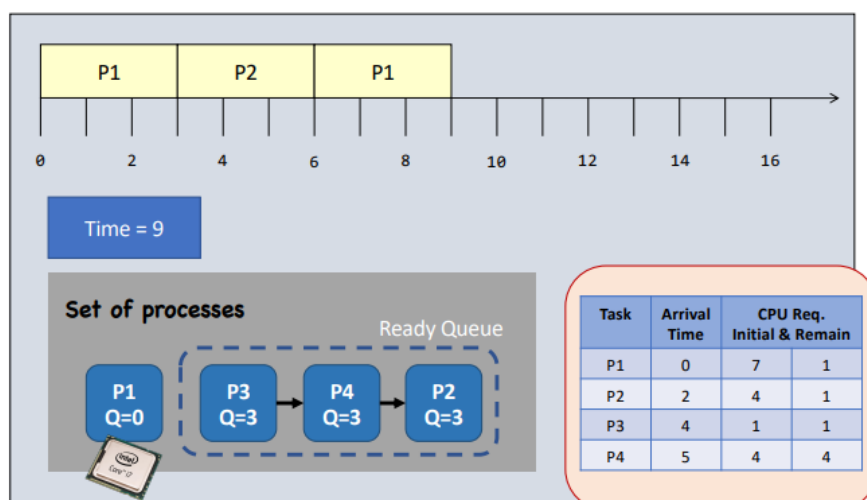
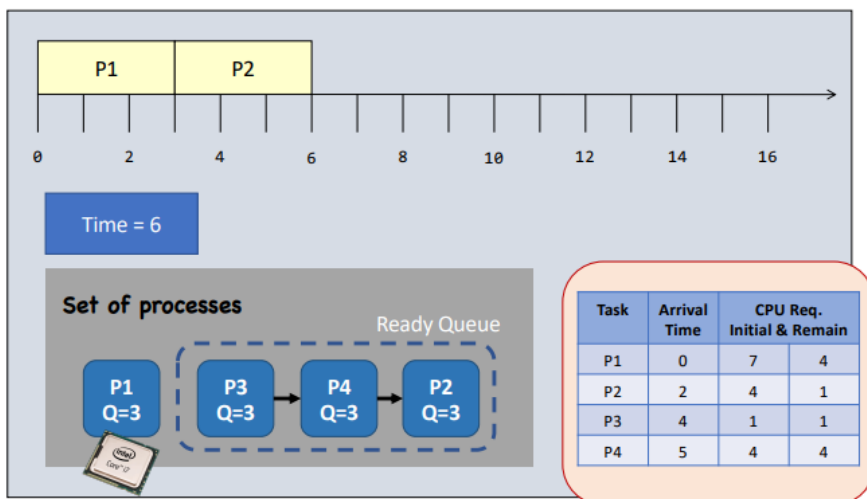
RR with Quantum

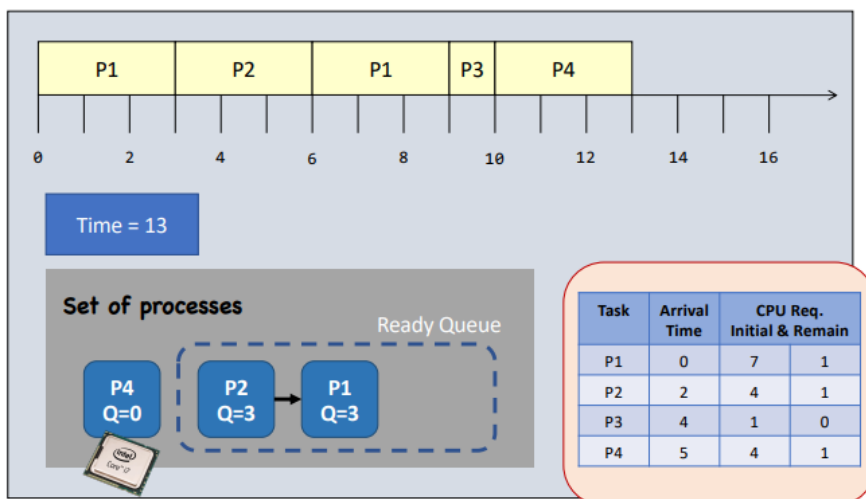
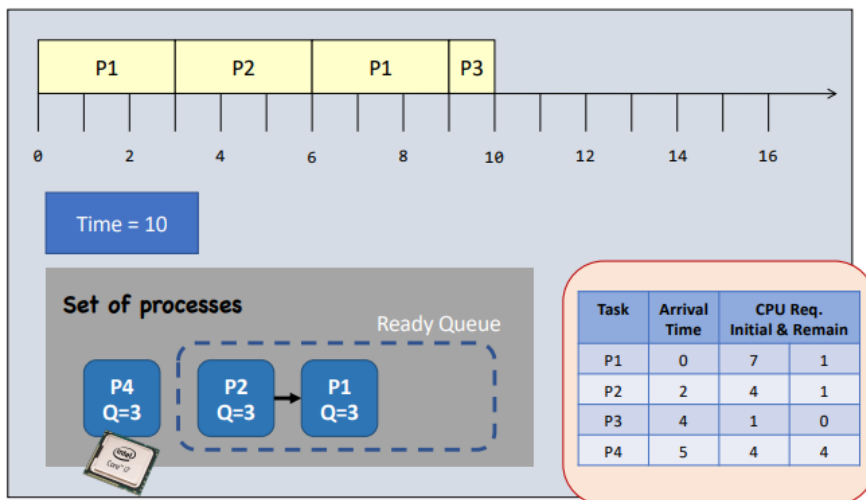
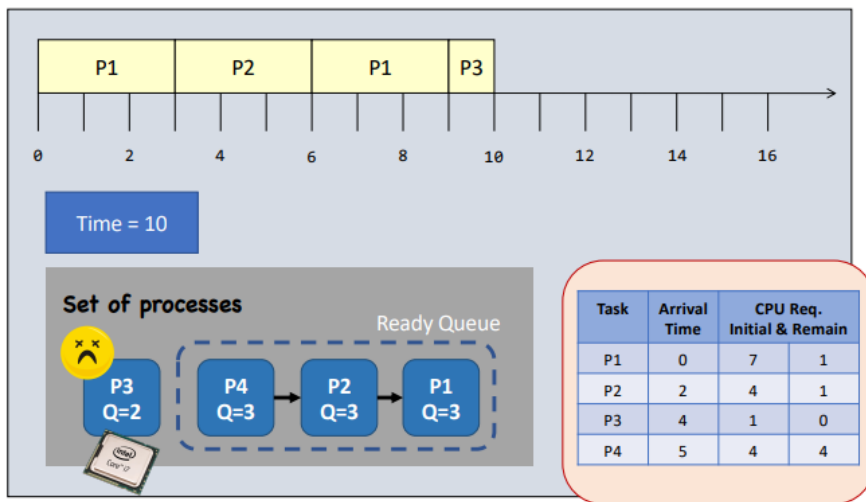
Suppose a RR schedule with Quantum = 3

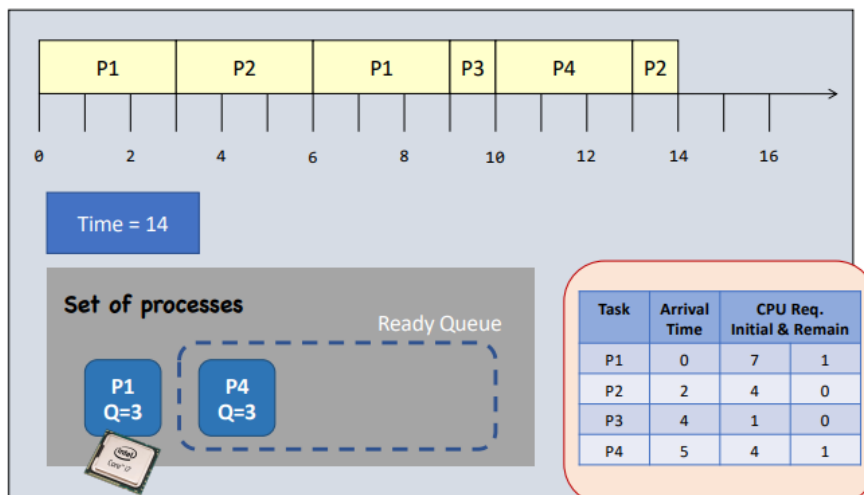
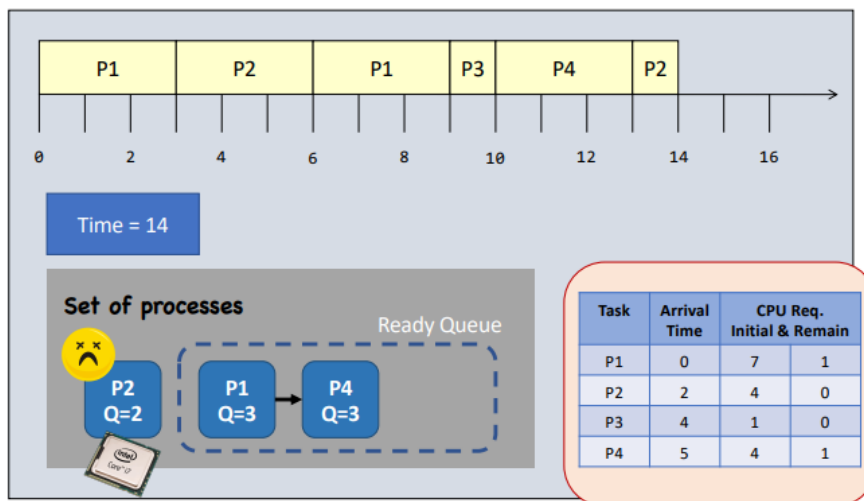
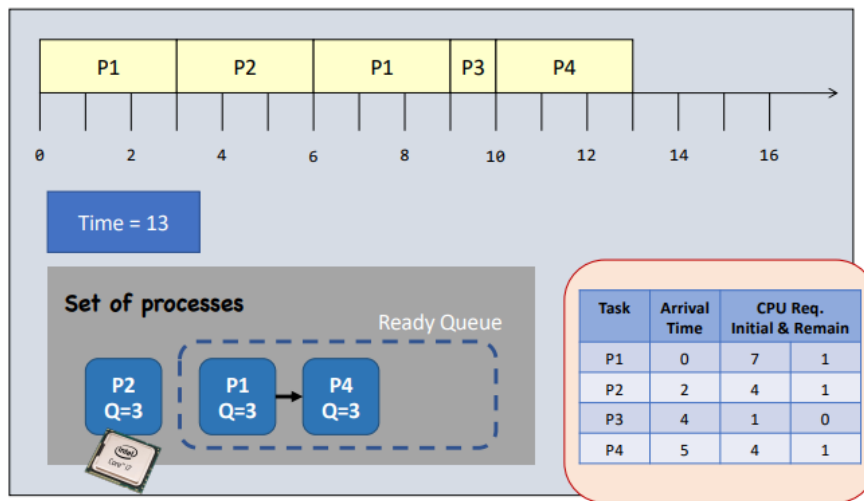


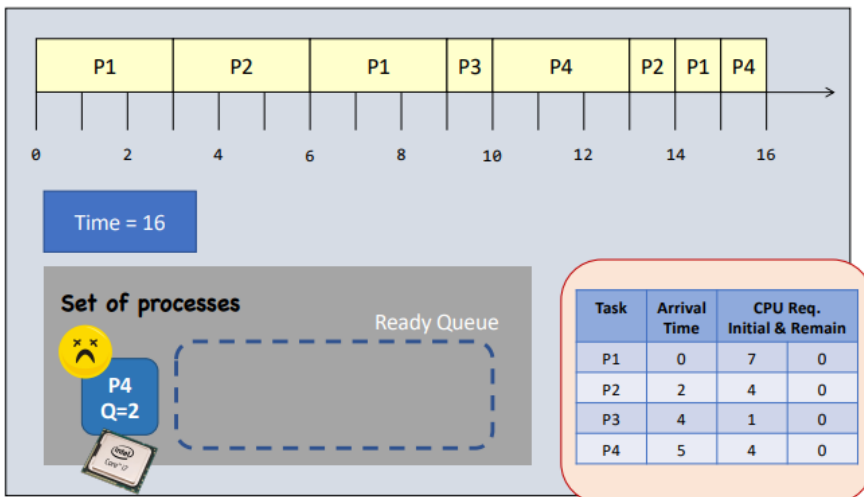
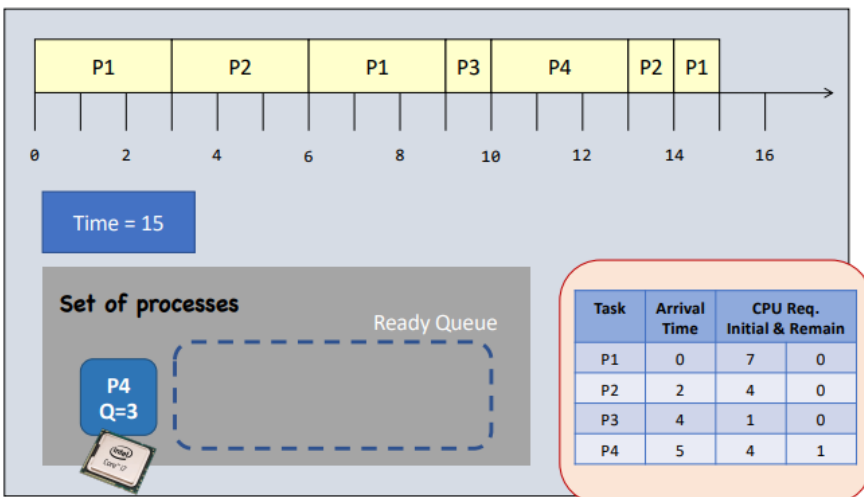
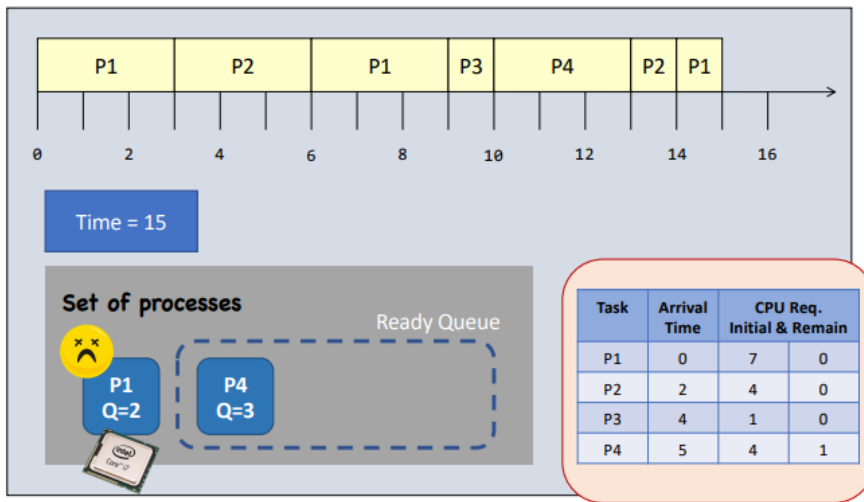












- waiting time
 - P1 = 8, P2 = 8, P3 = 5, P4 = 7
 - Average = $(8 + 8 + 5 + 7) / 4 = 7$
- turnaround time
 - P1 = 15, P2 = 12, P3 = 6, P4 = 11
 - Average = $(15 + 12 + 6 + 11) / 4 = 11$

RR v.s. SJF

	Non-preemptive SJF	Preemptive SJF	RR
Average waiting time	4	3	7 (largest)
Average turnaround time	8	7	11 (largest)
# of context switching	3	5	7 (largest)

So, the RR algorithm gets all the bad! Why do we still need it?

The responsiveness of the processes is great under the RR algorithm. E.g., you won't feel a job is "frozen" because every job gets the CPU from time to time!

Priority scheduling

- A priority number (integer) is associated with each process
- The CPU is allocated to the process with the highest priority (smallest integer = highest priority)
 - Non-preemptive: newly arrived process simply put into the queue
 - Preemptive: if the priority of the newly arrived process is higher than priority of the currently running process---preempt the CPU
- Static priority and dynamic priority
 - **static priority**: fixed priority throughout its lifetime
 - **dynamic priority**: priority changes over time

Problem

- Problem **Starvation**, low priority processes may never execute
- Solution: **Aging**, as time progresses increase the priority of the process
 - priority range from 127 (low) to 0 (high)
 - Increase priority of a waiting process by 1 every 15 minutes
 - 32 hours to reach priority 0 from 127

Completely Fair Scheduler

- Scheduling class
 - Default scheduling class: CFS
 - Real-time scheduling class
- Varying length scheduling quantum
 - Traditional UNIX scheduling uses 90ms fixed scheduling quantum
 - CFS assigns a proportion of CPU processing time to each task
- **Nice value**

- -20 to +19, default nice is 0
- **Lower nice value** indicates a **higher relative priority**
- Higher value is “being nice”
- Task with lower nice value receives higher proportion of CPU time
- **Virtual run time**
 - Each task has a per-task variable `vruntime`
 - Decay factor
 - Lower priority has higher rate of decay
 - nice = 0 virtual run time is identical to actual physical run time
 - A task with nice > 0 runs for 200 milliseconds, its `vruntime` will be **higher than** 200 milliseconds
 - A task with nice < 0 runs for 200 milliseconds, its `vruntime` will be **lower than** 200 milliseconds
- Lower virtual run time, higher priority
 - To decide which task to run next, scheduler chooses the task that has the **smallest** `vruntime` value
 - Higher priority can preempt lower priority

Example

- Two tasks have the same nice value
 - One task is I/O bound and the other is CPU bound
 - `vruntime` of I/O bound will be shorter than `vruntime` of CPU bound
 - I/O bound task will eventually have higher priority and preempt CPU-bound tasks whenever it is ready to run