

# Lecture8 Demand Paging

## 1. Demand Paging Mechanisms


### How to go beyond physical memory

- How to support large address space?
  - 64-bit machine supports up to 4EB address space
  - Applications may use more space than available in physical memory
- Solution: **stash away portions of address spaces** that aren't currently in use
  - in the **next-level of storage** (e.g., hard disk drive)
  - slower but much larger

### An abstraction of Address Space


- **Application:** memory overlays
  - Application in charge of moving data between memory and disk
  - e.g., calling a function needs to make sure the code is in memory
- **OS:** demand paging
  - OS **configures page table entries**
  - Virtual page maps to physical memory or files in disk
  - **Process** sees an **abstraction** of address space
  - OS determines **where the data is stored**

### Swap Space



- Swap space is a partition or a file stored on the **disk**
  - OS swaps pages out of memory to it
  - OS swaps pages from it into memory
- Swap space conceptually divided into **page-sized units**
  - OS maintains a **disk address** of each page-sized unit

## Example



- 4-page physical memory and an 8-page swap space
  - Proc 0 has three virtual pages
  - Proc 1 has four virtual pages
  - Proc 2 and Proc 3 each has two virtual pages

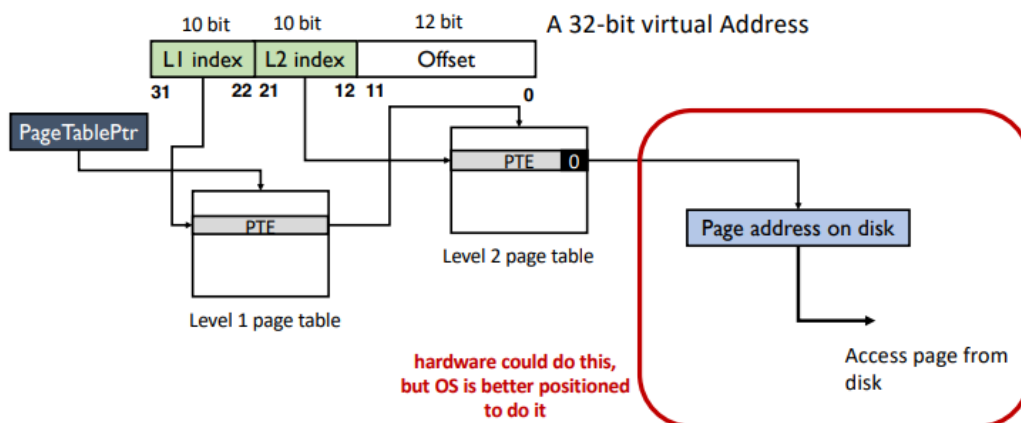
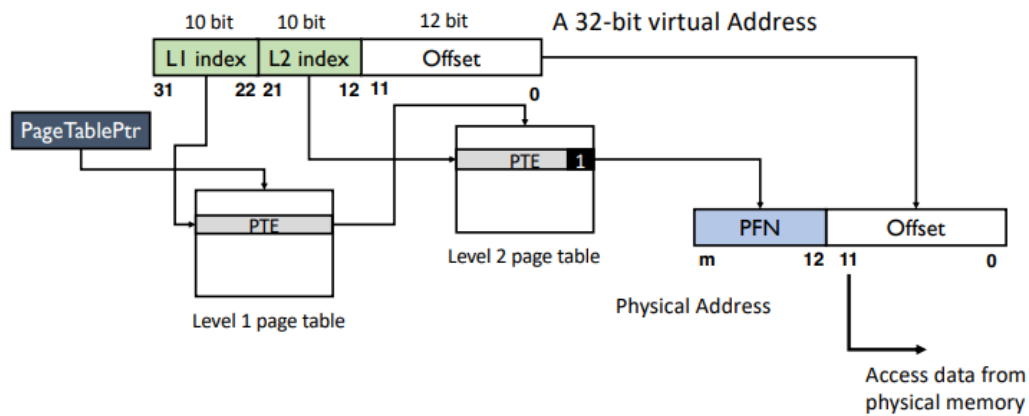
## Demand Paging

- Load pages from disk to memory only as they are needed
  - Pages are loaded “on demand”
- Data transferred in the **unit of pages**
- Two possible on-disk locations
  - **Swap space**
    - created by OS for temporary storage of pages on disk
    - e.g., pages for stack and heap
  - **Program binary files**
    - The code pages from this binary are only loaded into memory when they are executed
    - Read-only, thus never write back

## Physical Memory as a Cache

- Physical memory can be regarded as a cache of on-disk swap space
- Block size of the cache: 1 page(4 KB)
- Cache organization (direct-mapped, set-associative, fully-associative)
  - Fully associative: any disk page maps to any page frame
- Page replacement policy
  - LRU, Random, FIFO
- Page miss
  - Go to lower level to fill page (i.e. disk)

## Present Bit



## Page Faults

- Present bit = 0 raises a page fault exception
  - OS gets involved in address translation
- Page fault handler
  - Find **free page frame** in **physical memory**
    - Find one free page frame from a free-page list
    - If no free page, trigger **page replacement**
  - Fetch page from **disk** and store it in physical memory
    - Determine the faulting **virtual address** from register
    - Locate the **disk address** of the page in PTE (where PFN should be stored)
    - Issues a request to disk to fetch the page into memory
    - Wait ..... (could be a very long time, **context switch**)
    - When I/O completes, update page table entry, PFN, present bit
- After page fault
  - Return from page fault **exception**
  - CPU **re-execute the instruction** that accesses the virtual memory
  - No more page fault since **present bit is set** this time

- TLB entry loaded from PTE

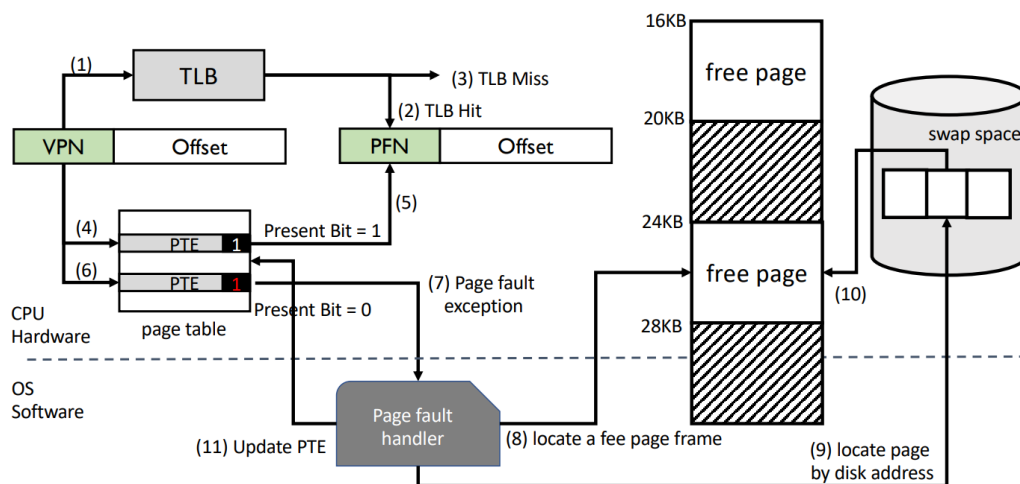
## Page Replacement

- find a page frame to be replaced
  - Page replacement policy decides which one to replace
- If page frame to be replaced is dirty, write it back to disk
- Update all PTEs pointing to the page frame
- Invalidate all TLB entries for these PTEs

When to trigger page replacement

- [Lecture13 Deadlock.assets](#) Proactive page replacement usually leads to better performance
  - Page replacement even though no one needs free page frames (yet)
  - Always reserve some free page frames in the system
- Swap daemon
  - background process for reclaiming page frames
  - **Low watermark**: a threshold to trigger swap daemon
  - **High watermark**: a threshold to stop reclaiming page frames

## Put it all together



## 2. Page Replacement Policy

### Effective Access Time (EAT)

- $EAT = \text{Hit Rate} \times \text{Hit Time} + \text{Miss Rate} \times \text{Miss Penalty}$
- Example
  - Memory access time: 200ns
  - Average page-fault service time: 8 ms
  - Suppose  $p$  = Probability of miss,  $1 - p$  = Probability of hit

- $EAT = (1-p) \times 200ns + p \times 8ms$
- If one access out of 1,000 causes a page fault, then EAT is about  $8.2 \mu s$ 
  - This is a slowdown by a factor of 40
- If we want slowdown by less than 10%
  - which means  $EAT < 220 ms$
  - This is about 1 page fault in 400,000

## Types of Cache Misses

- Compulsory Misses
  - **Cold-start miss:** pages that have **never** been fetched into memory before
  - Prefetching: loading them into memory before needed
- Capacity Misses
  - **Not enough memory:** must somehow increase available memory size
  - One option: Increase amount of DRAM (not quick fix!)
  - Another option: If **multiple** processes in memory: adjust percentage of memory allocated to each one
- Conflict Misses
  - fully-associative cache (OS page cache) does not have conflict misses

## Page Replacement Policies

Suppose we have 3 page frames, 4 virtual pages, and following reference string:

- A B C A B D A D B C B

### Optimal / MIN

- Replace page that will not be used for the longest time
- Lead to minimum page faults in theory

Ref:	A	B	C	A	B	D	A	D	B	C	B
Page:											
1	A									C	
2		B									
3			C			D					

- MIN: 5 faults

## FIFO

- Throw out oldest page first
- May throw out heavily used pages instead of infrequently used

Ref:	A	B	C	A	B	D	A	D	B	C	B
Page:											
1	A					D				C	
2		B					A				
3			C						B		

## RANDOM

- Pick random page for every replacement
- Pretty unpredictable – makes it hard to make real-time guarantees

## Least Recently Used (LRU)

- Replace page that has not been used for the longest time
- **Temporal locality** of program
  - If a page has not been used for a while, it is unlikely to be used in the near future

Ref:	A	B	C	A	B	D	A	D	B	C	B
Page:											
1	A									C	
2		B									
3			C			D					


- LRU is not always the optimal
  - Consider the following reference string: A B C D A B C D A B C D

Ref:	A	B	C	D	A	B	C	D	A	B	C	D
Page:												
1	A			D			C			B		
2		B			A			D			C	
3			C			B			A			D

## Least Frequently Used (LFU)

- Replace page that has not been accessed many times
- **Spatial locality** of program
  - If a page has been accessed many times, perhaps it should not be replaced as it clearly has some value

## Bélády's Anomaly



- One desirable property: When you add memory the miss rate drops
  - Yes for **LRU** and **MIN**
  - Not necessarily for **FIFO**. For **FIFO**, more page frames may lead to more page faults

## Example


Ref: Page:	A	B	C	D	A	B	E	A	B	C	D	E
1	A			D			E					
2		B			A					C		
3			C			B					D	

Ref: Page:	A	B	C	D	A	B	E	A	B	C	D	E
1	A						E				D	
2		B						A				E
3			C						B			
4				D						C		

## Page Fault Curve

Page fault curve for FIFO on reference string

- 70120304230321201701



### 3. LRU Implementation

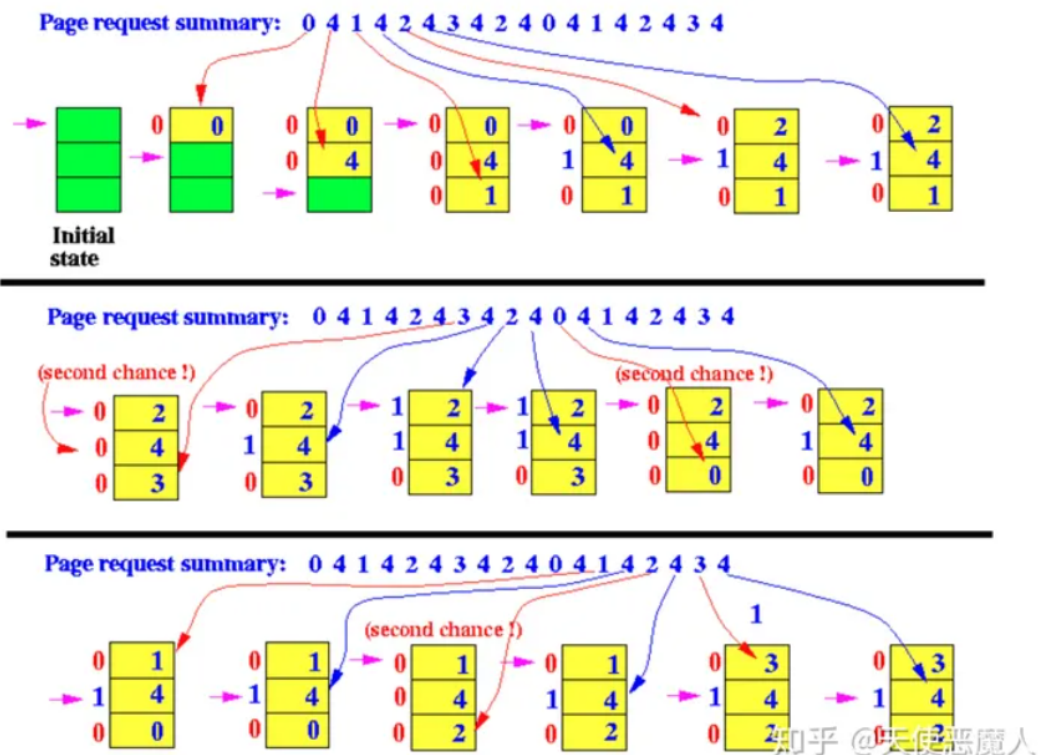
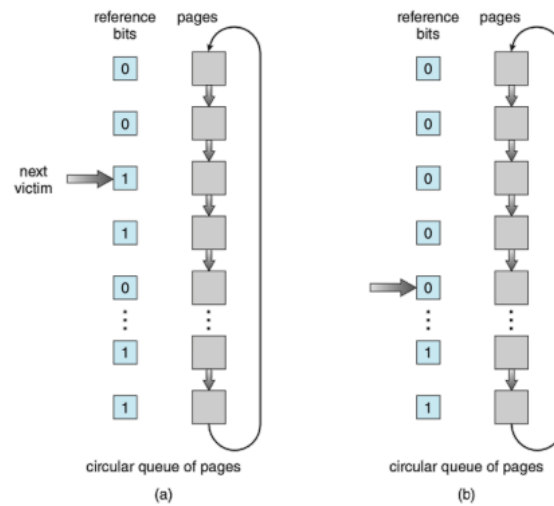
- Hardware support is necessary
  - Update a data structure in OS upon every memory access
  - E.g., a timestamp counter for each page frame
- Overhead
  - One **additional memory write** for each memory access
    - TLB hit does not save the extra memory access
  - Scan the entire memory to find the LRU one
    - 4GB physical memory has 1 million page frames
    - sorting is time consuming

### LRU Approximation with Reference Bit

- Reference bit
  - **One reference bit** per page frame
  - All bits are cleared to **0** initially
  - The first time a page is referenced, the reference bit is **set by CPU**
    - Can be integrate with page table walk
  - The order of page accesses approximated by two clusters: **used** and **unused** pages
- Examples
  - Clock algorithm (also called second-chance algorithm)
  - Enhanced clock algorithm with dirty bits

### Clock Algorithm





- Arrange physical pages in a circular list
- CPU sets reference bit to 1 upon first access
- OS maintains a **pointer**
  - When a replacement occur, check reference bit of the current page
  - **If 1:** the page has been accessed recently, clear the bit (set to 0) and move to the next page
  - **If 0:** the page has not been accessed recently, good candidate for replacement, stop

## With dirty bit

- Enhance clock algorithm with a dirty bit
  - dirty bit = 1: the page **has recently been modified**
- CPU sets dirty bit to 1 upon **write** access


- When a replacement occurs, OS checks (ref bit, dirty bit), and selects a candidate page in decreasing order

Reference bit	Dirty bit	Description
0	0	neither recently used nor modified — best page to replace
0	1	not recently used but modified — not quite as good, because the page will need to be written out before replacement
1	0	recently used but clean — probably will be used again soon
1	1	recently used and modified — probably will be used again soon, and the page will be need to be written out to secondary storage before it can be replaced

## LRU Approximation with Reference Bit and Counter

- Each physical page frame is associated with one **reference bit** and a **counter**
  - **Reference bit** indicate recent access
    - set by CPU hardware, cleared by OS
  - **Counter** records history of accesses
    - Maintained by OS

## Additional-reference-bits Algorithm



- 8-bit history register associated with each page frame
- **Timer interrupt** every 100ms
  - reference bit shifts to **highest bit** in the history register
  - other bits shift right and discard the lowest bit
  - 00000000 unused page in 800ms
- Compare history register as unsigned integer
  - Larger value more recently used
  - 11000100 > 01110111

- Approximate LRU with more bits and more frequent interrupts

## Nth-chance Clock Algorithm

- All page frames arranged in a **circular list** and each page frame is associated with a reference bit and a counter
- CPU hardware sets reference bit upon **memory accesses**
- OS checks the reference bit of the page pointed to by the clock hand
  - 1 -> clear reference bit and the counter
  - 0 -> increment counter; if count = N, replace page
- How do we pick N
  - Large N: better approximation of LRU
    - If N = 1K, really good approximation
  - Small N: more efficient
    - otherwise might have to look a long way to find free page

## 4. Page Frame Allocation

- How do we allocate memory among different processes?
  - Does every process get the same fraction of memory? Different fractions
  - Should we completely swap some processes out of memory?
- Minimum number of pages per process
  - Depends on the **computer architecture**
  - How many pages would one instruction use at one time
    - x86 only allows data movement between memory and register and no indirect reference
    - needs at least one instruction page, one data page, and some page table pages
- Maximum number of pages per process
  - Depends on available physical memory

## Global versus Local Allocation

- Global replacement
  - Process selects replacement frame from **all page frames**
  - One process can take a frame from another process
- Local replacement
  - Each process selects from only its **own set of allocated frames**

## Allocation Algorithms

- Equal allocation
  - Every process gets same amount of memory
  - Example: 100 frames, 5 processes -> process gets 20 frames
- Proportional allocation
  - Number of page frames proportional to the size of process
  - $s_i$  = size of process  $p_i$  and  $m$  = total number of frame
  - $a_i$  = allocation for  $p_i = m \times \frac{s_i}{\sum s_j}$
- Priority allocation
  - Number of page frames proportional to the priority of process
  - Possible behavior: If process  $p_i$  generates a page fault, select for **replacement a frame from a process with lower priority number**

## Trashing

- The **memory demands** of the set of running processes simply **exceeds** the **available physical memory**
- Early OS
  - Working set: the pages used actively of a process
  - Reduce the # of process so their working set fit into memory
- Modern OS
  - Out-of-memory killer when memory is oversubscribed
  - May need a reboot